

# IOM AI School

TOPIC: Deep Learning – Theory and Practice  
SPEAKER: Matthias Täschner

Using materials from Laura Zigutyte, Michaela Unger (EKfZ, TU Dresden), Robert Haase (ScaDS.AI, Leipzig University)  
These slides can be reused under the terms of the [CC-BY4.0](https://creativecommons.org/licenses/by/4.0/) license.

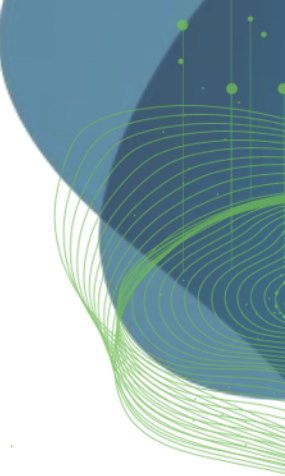


**SACHSEN** Diese Maßnahme wird gefördert durch die Bundesregierung aufgrund eines Beschlusses des Deutschen Bundestages. Diese Maßnahme wird mitfinanziert durch Steuermittel auf der Grundlage des von den Abgeordneten des Sächsischen Landtags beschlossenen Haushaltes.



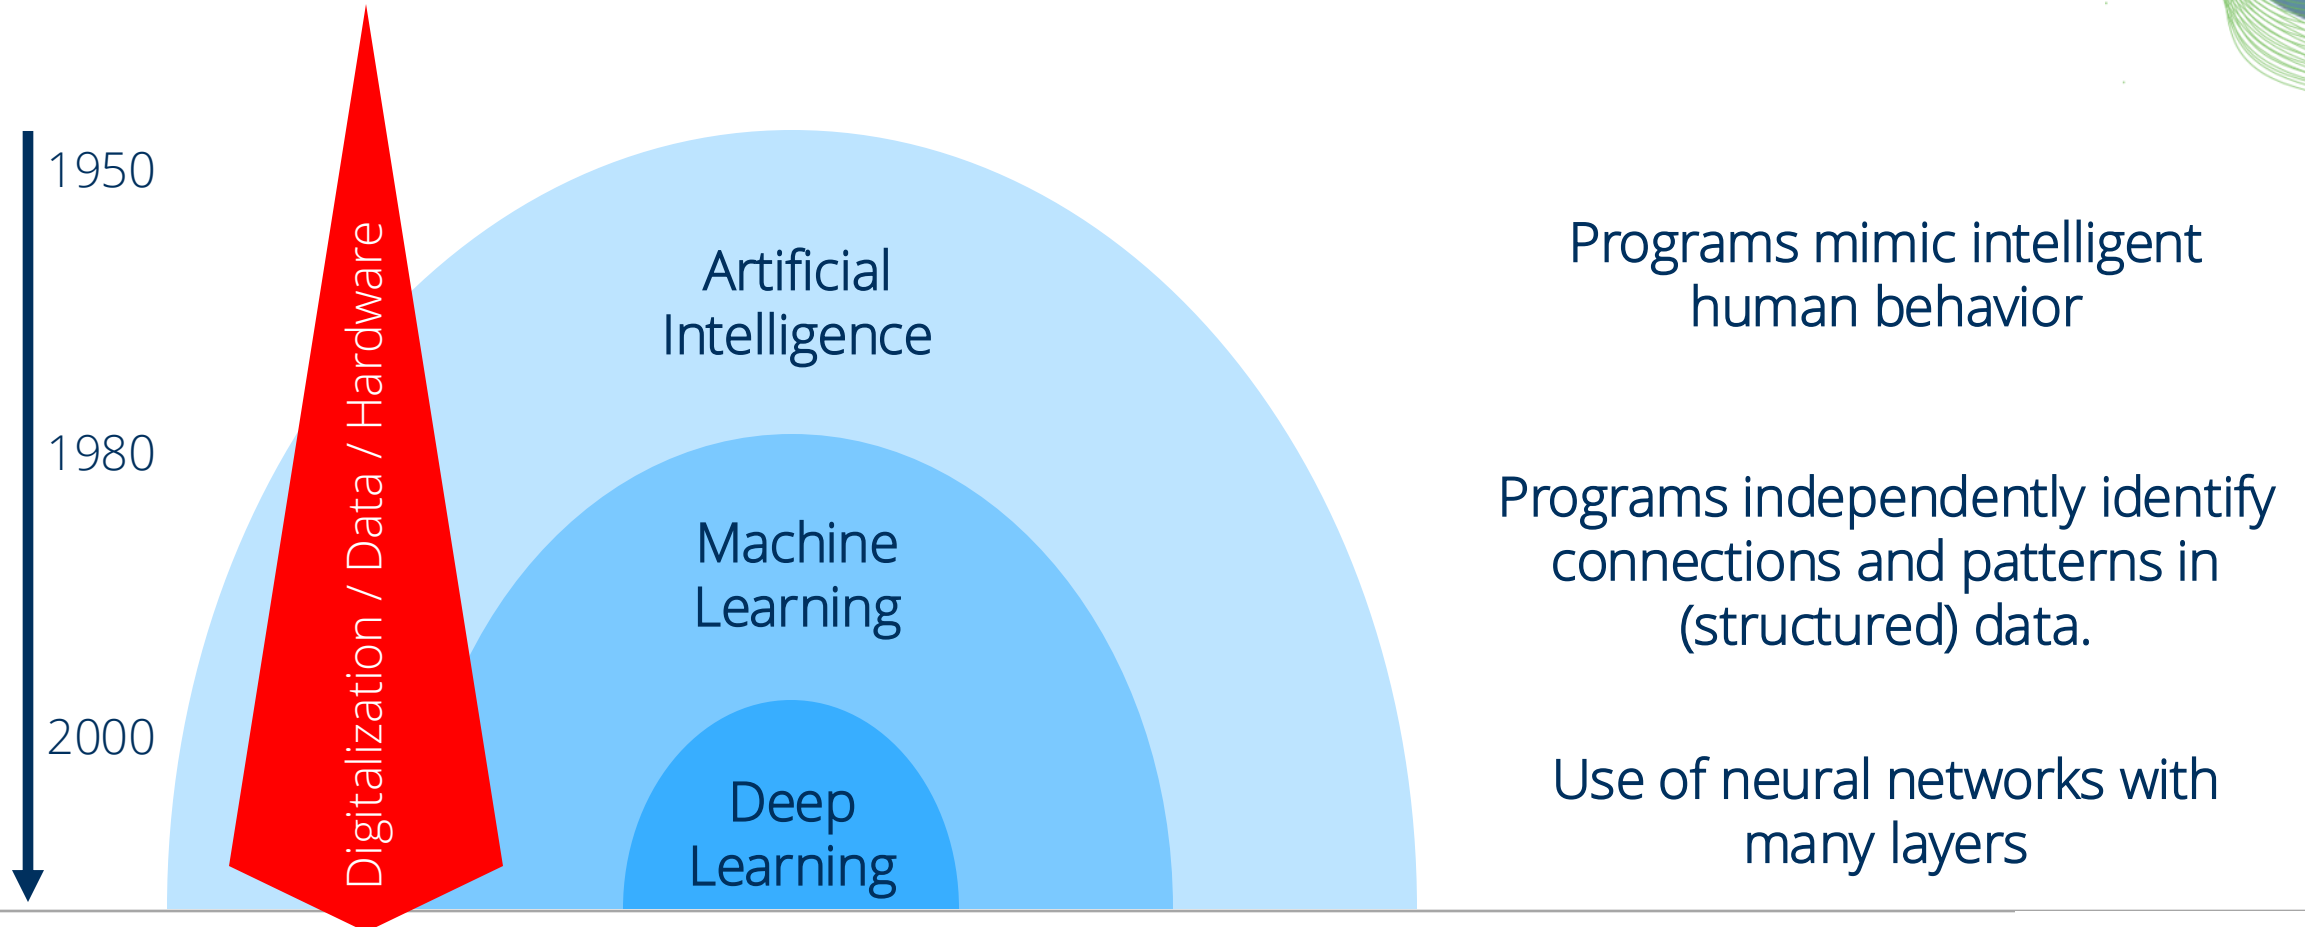
# AGENDA

- Areas of Artificial Intelligence (AI)
- Neural Networks and Deep Learning
- Neural Network Architectures
- Practice: Neural Network for Regression



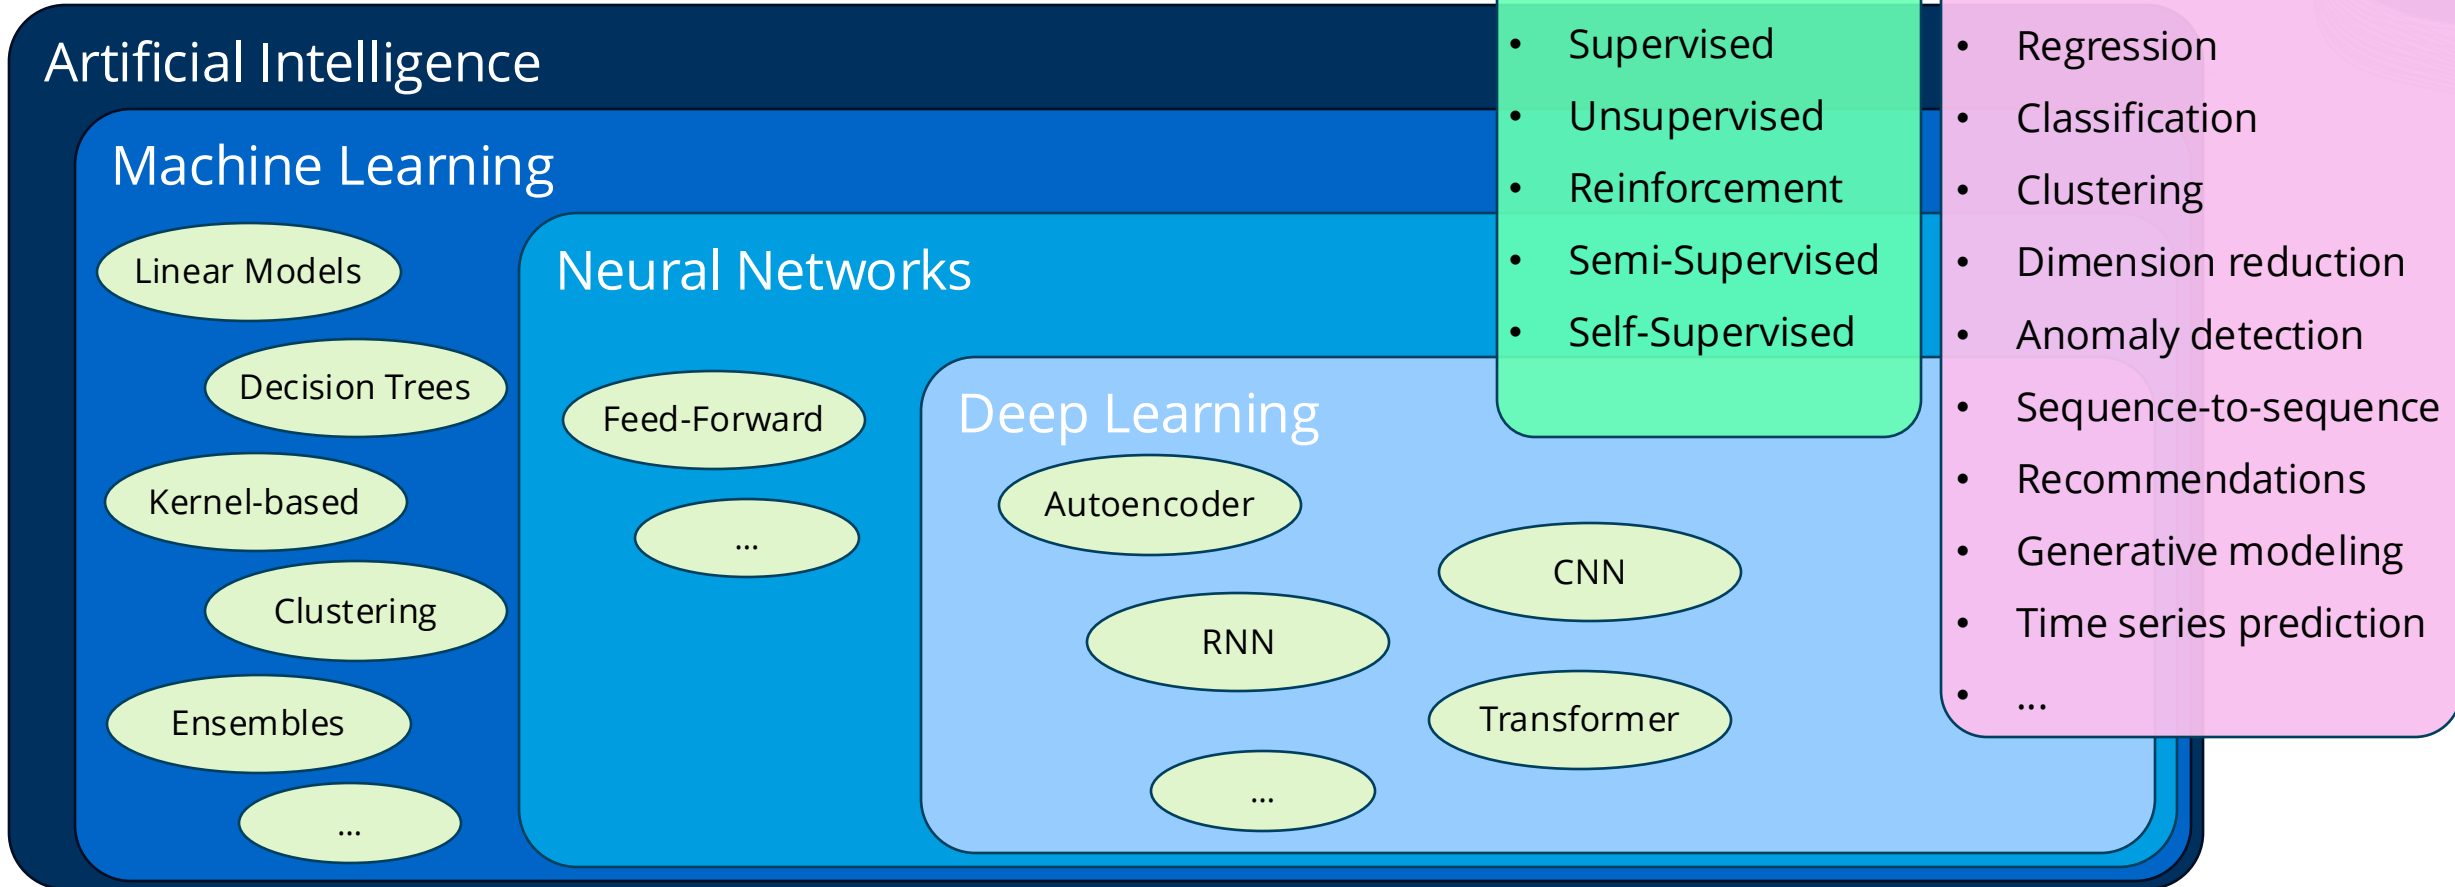
# Areas of Artificial Intelligence

Historical phases of AI



# Areas of Artificial Intelligence

Areas, paradigms, model families (not exhaustive)



# Areas of Artificial Intelligence

## Paradigms - Supervised Learning

### Procedure

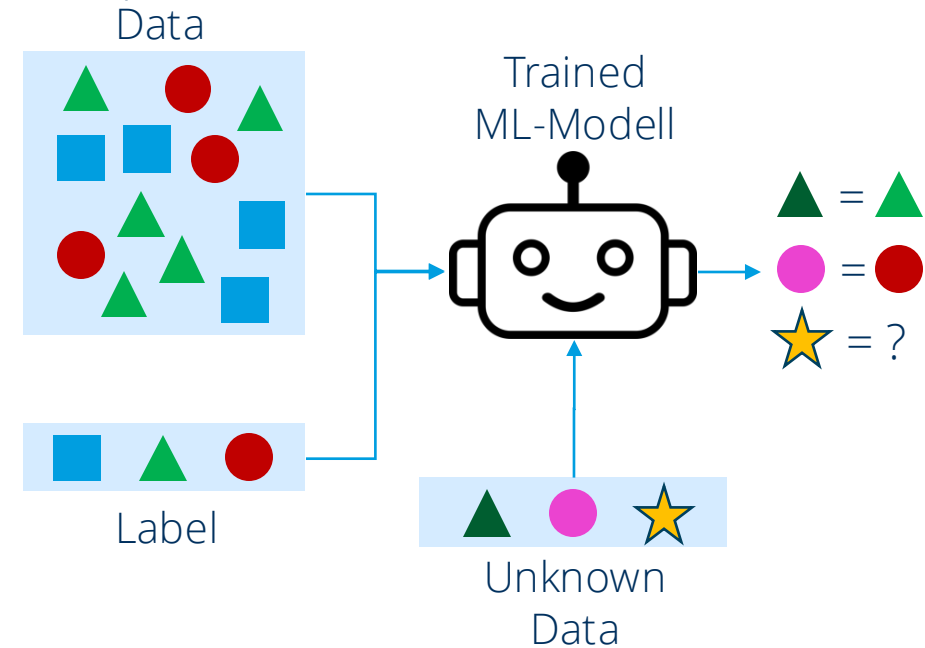
- ML models are trained using pre-labeled data
- Input data and the desired target values (ground truth) are provided for training
- Prediction of target values on new, previously unknown input data of the same format

### Application examples

- Classification, Regression
- Anomaly Detection
- Generative modeling

### Algorithm / model type examples

- Linear Regression
- Decision Trees & Random Forest (DT & RF)
- Support-Vector-Machines (SVM)
- Neural Networks (NN)



# Areas of Artificial Intelligence Paradigms - Unsupervised Learning

## Procedure

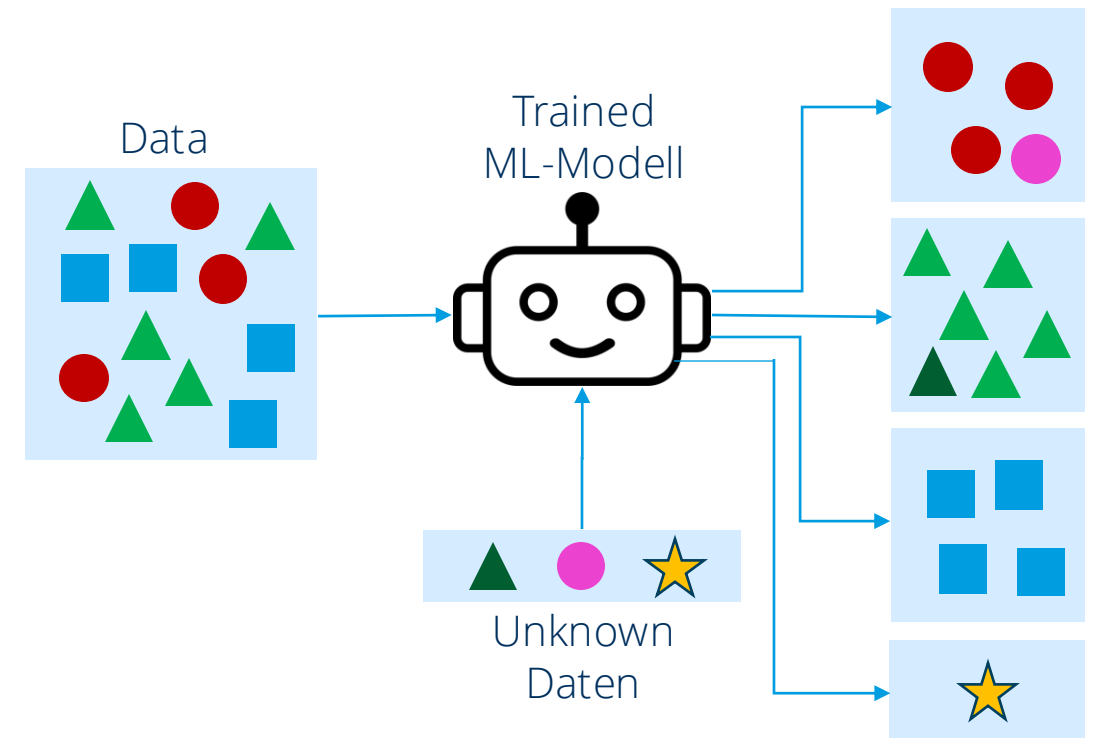
- ML models are trained with unlabeled data
- Models independently recognize patterns, structures, and correlations

## Application examples

- Clustering
- Dimensionality Reduction
- Anomaly Detection

## Algorithm / model type examples

- K-Means Clustering
- Principal Component Analysis (PCA)
- Autoencoder



# Areas of Artificial Intelligence Paradigms - Reinforcement Learning

## Procedure

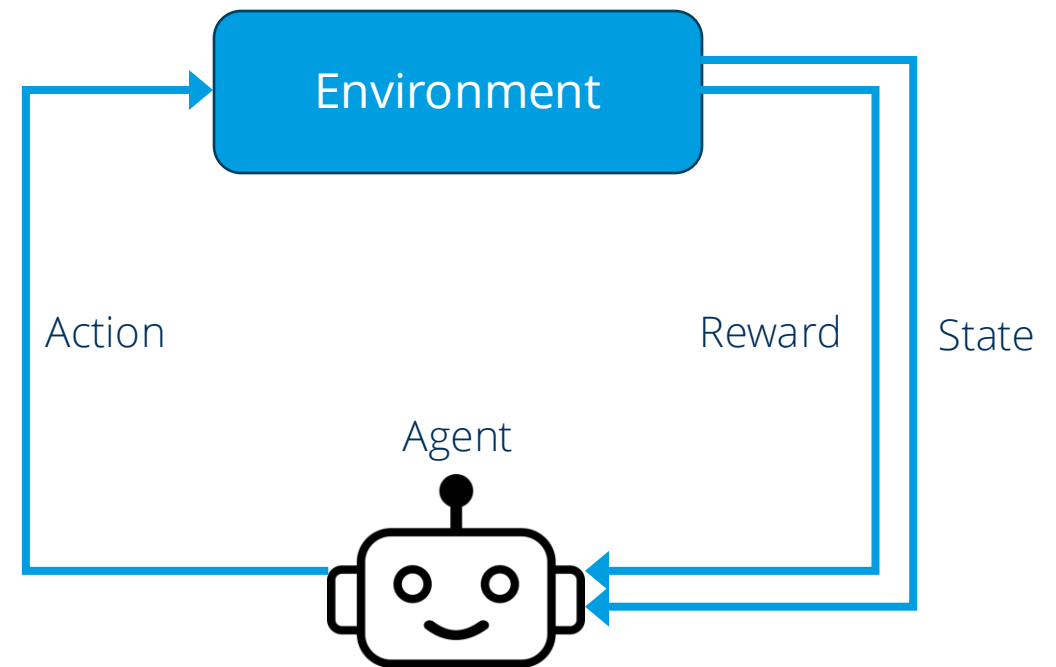
- An agent is trained to maximize a specific reward in an environment through trial and error
- Rules define the possible actions the agent can take
- Rewards (or punishments) influence the agent's behavior

## Application examples

- Game agents, e.g., in chess or Go
- Automation systems, robotics
- Simulations

## Algorithm / model type examples

- Q-Learning
- Markov Decision Processes (MDP)
- Monte Carlo Methods






# Areas of Artificial Intelligence Paradigms - Mixed Forms

## Semi-Supervised

- A small dataset with labeled samples
- A larger dataset with unlabeled samples from the same distribution
- Training on labeled data → model assigns “pseudo-labels” to unlabeled data
- Include data with the most reliable “pseudo-labels” in further training (human-in-the-loop)

## Self-Supervised

- Data is not labeled; required labels are generated by the model itself
- The model generates its own “practice questions” and answers from the raw data
- In doing so, it learns useful patterns from the data
- Example: Cover words in a sentence and then guess them – no label necessary





# Areas of Artificial Intelligence

## Linear models vs. non-linear models

- Linear model  $\neq$  plotting a straight line / curve
- Linearity refers to model parameters


*A model is linear if its prediction is a linear combination of the features.  
(e.g., weighted sum of input data)*

- Linear models
  - Easy to interpret, based on well-understood principles, fast and efficient training
  - Can only model simple, linear relationships in the data
  - Examples: Linear or logistic regression, ARIMA, linear SVM, etc.
- Non-linear models
  - Allows to model complex, non-linear relationships in the data
  - More difficult to interpret, complex training with more parameters, often more data required
  - Examples: Decision tree, kernel SVM, k-NN, neural networks, etc.

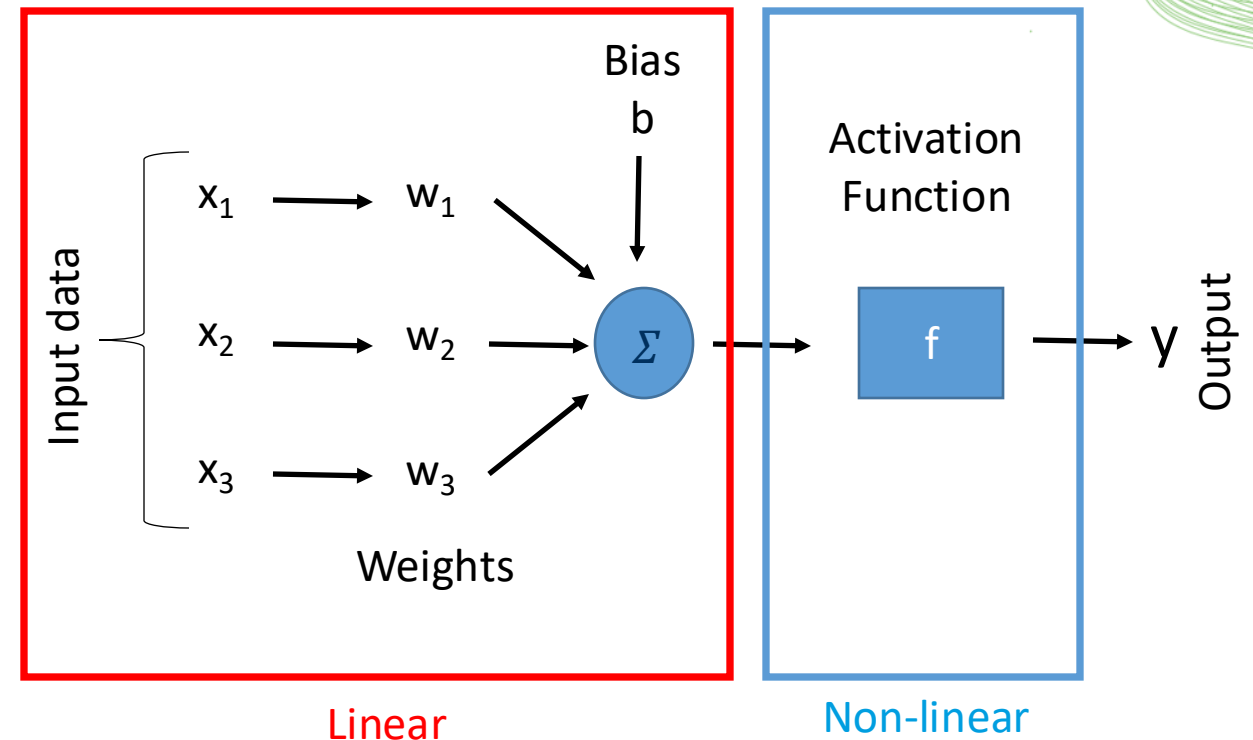
# Neural Networks and Deep Learning

## Basic Architecture

Neurons from a biological perspective



Neuron in Data Science



$$y = f(w_1x_1 + w_2x_2 + w_3x_3 + b)$$

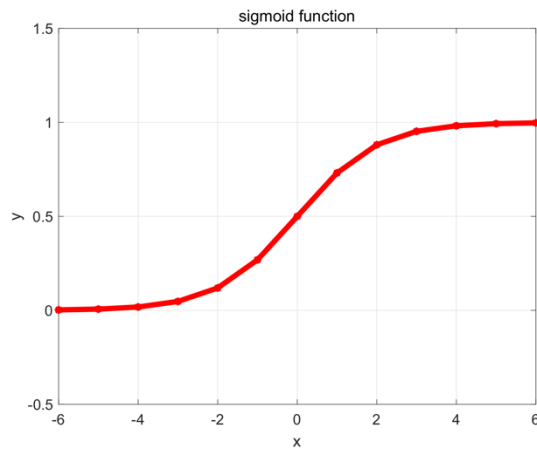
Source: [https://commons.wikimedia.org/wiki/File:Blausen\\_0657\\_MultipolarNeuron.png](https://commons.wikimedia.org/wiki/File:Blausen_0657_MultipolarNeuron.png)  
License: [CC-BY 3.0, BruceBlas](#)

# Neural Networks and Deep Learning

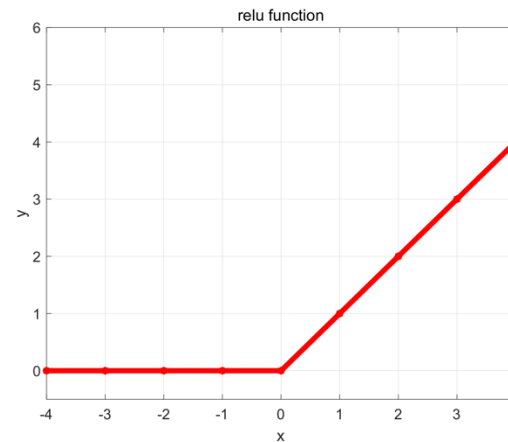
## Basic Architecture

### Activation-Function

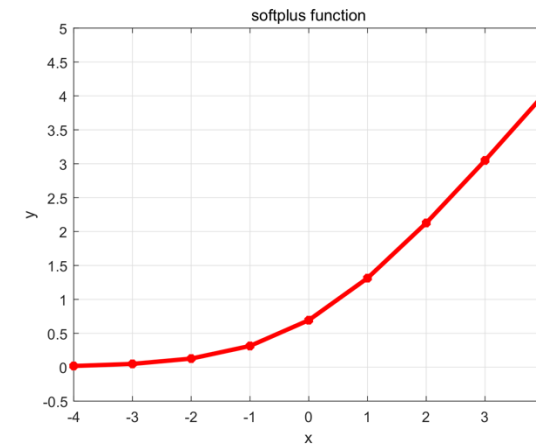
- Non-linear transformation of the linear combination of input features



(a) The curve of sigmoid function



(c) The curve of ReLU function



(d) The curve of softplus function

Source: Wang, Y., Li, Y., Song, Y., & Rong, X. (2020). The Influence of the Activation Function in a Convolution Neural Network Model of Facial Expression Recognition. Applied Sciences, 10(5), 1897. <https://doi.org/10.3390/app10051897>, OpenAccess

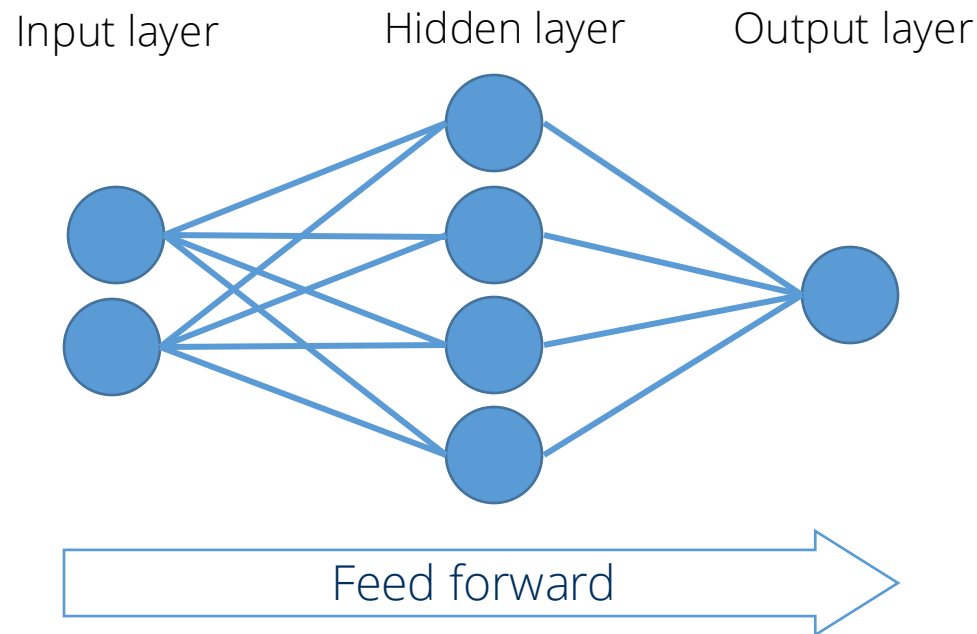
- More examples and explanations: [https://en.wikipedia.org/wiki/Activation\\_function](https://en.wikipedia.org/wiki/Activation_function)

# Neural Networks and Deep Learning

## Basic Architecture

### Multi-Layer Perceptron

- Several layers of neurons: Input → Hidden → Output
- Neurons in one layer are fully connected to those in the next layer (Fully-Connected, Dense)
- Information flows through the network from input to output only (Feed Forward)



**„Deep Learning“**  
Neural networks with  
**many** hidden layers

# Neural Networks and Deep Learning


## Basic Architecture

Data input and information flow  
(Fully-Connected Feed-Forward)

Year	m <sup>2</sup>	Zip-Code	Price
1991	234	12345	xxxxx
1896	400	54638	xxxxx
2005	185	95463	xxxxx




$$y = f(w_1x_1 + w_2x_2 + w_3x_3 + b)$$



■ Weight

■ Bias



# Neural Networks and Deep Learning


## Data Processing

### Processing and conversion of input data

- Create features that can be used and learned by ML models / neural networks
- Goal: Turn input data into numerical representation so a neural network can process them



Is the data representative?  
Handling of missing data  
Handling of outliers  
Feature selection  
Data scaling / normalization  
Embeddings  
Train-test-split  
...




# Neural Networks and Deep Learning

## Training Process

### Learning through “Back Propagation”

- Step 1: Initialize weights of all neurons with small random values
- Step 2: Send data point (**sample**) through the network for prediction (**forward pass**)
- Step 3: Compare prediction with known label for the sample (supervised learning with **ground truth**), calculate error between prediction and ground truth (via **loss function**)
- Step 4: Starting with the output layer (**backward pass**), adjust the weights in the neurons with the goal to minimize the calculated error (e.g., via **gradient descent**)
- Repeat steps 2-4 for each sample in training data (one **epoch**), running through this for several epochs



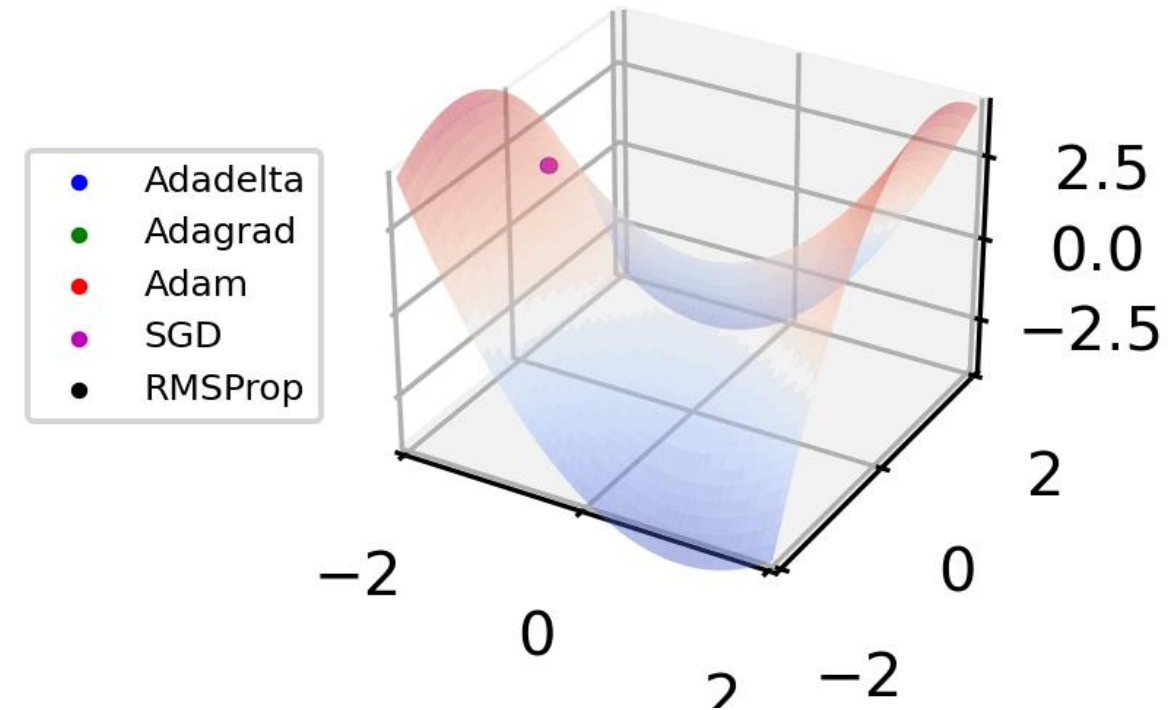


# Neural Networks and Deep Learning

## Training Process

### Weight Adjustment – Optimization via Gradient Descent

- Gradient descent updates weights to minimize an objective function → loss function → error
- Optimizer follows the steepest descent (slope) on the loss function
- Learning rate controls the step size
  - too large = overshoot
  - too small = slow training
- Training may get stuck in local minima instead of reaching best global solution
- Different optimizer algorithms improve training



Source: VirtualVistas - Own work,  
<https://commons.wikimedia.org/w/index.php?curid=139889895>, CC BY-SA 4.0

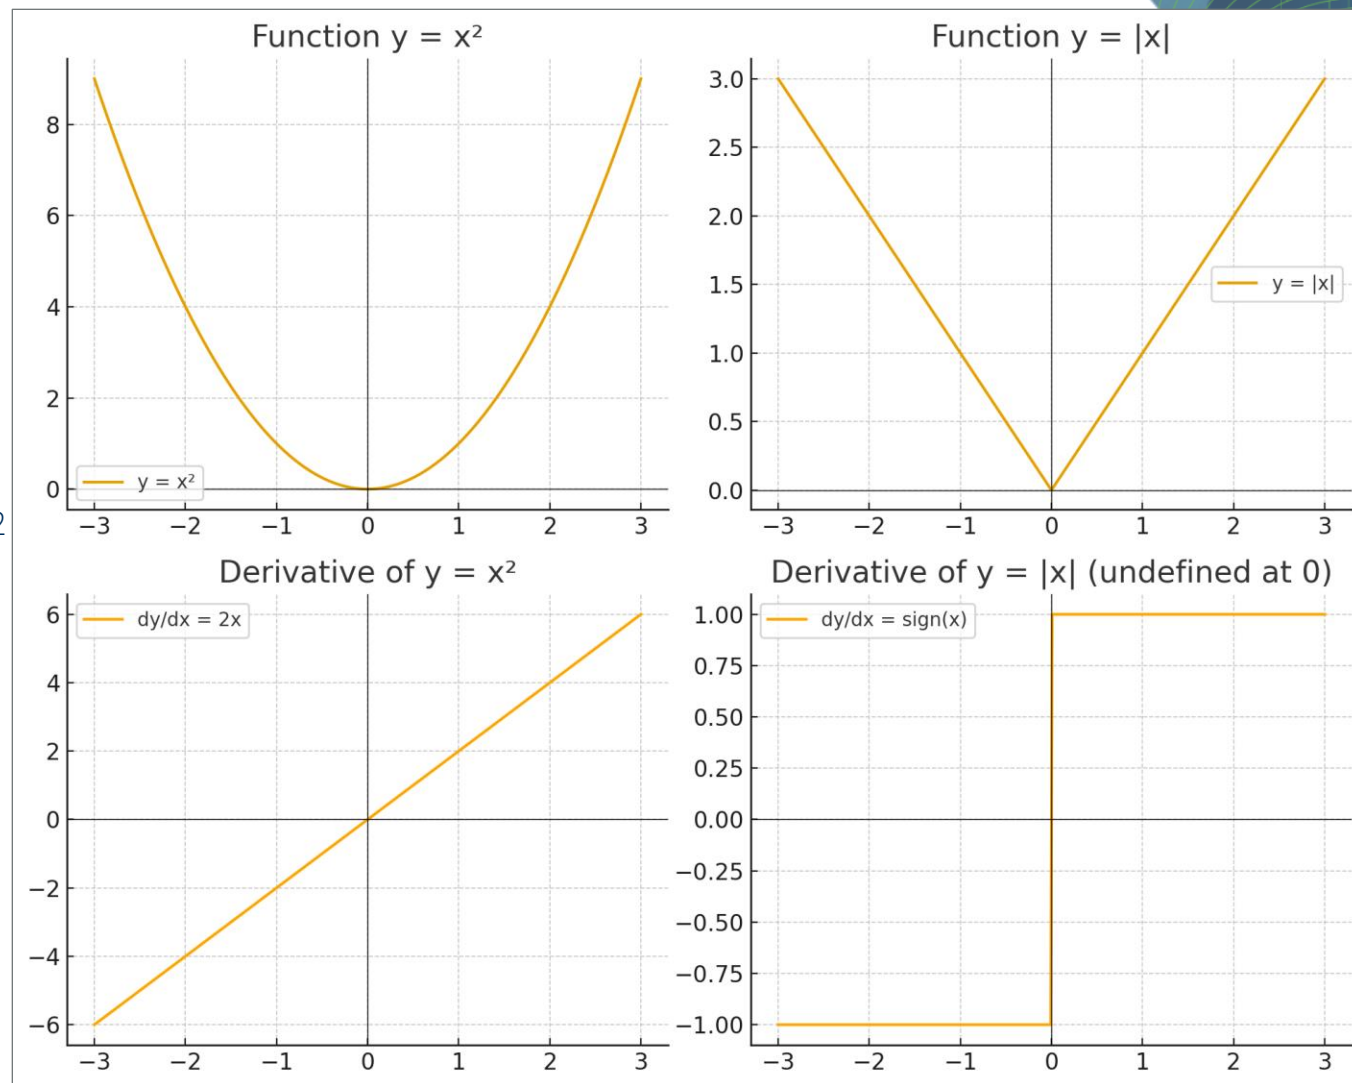


# Neural Networks and Deep Learning

## Training Process

### Loss Function

- Needs to be a differentiable function to compute a slope
- Derivate must be defined for all points
- Example:
  - Mean Squared Error:  $(y_{\text{true}} - y_{\text{predicted}})^2$



# Neural Networks and Deep Learning

## Training Process

### Architecture and Training Parameters (Hyperparameters)


- Model parameters = learned by the neural network during training (weights, biases)
- Model architecture choices
  - Layer types and arrangement
  - Number of neurons per layer
  - Type of activation function
  - ...
- Training parameter
  - Number of samples / batches per epoch
  - Number of epochs
  - Optimization method and learning rate
  - ...
- Hyperparameter optimization → find the best configuration for the architecture and the training

# Neural Networks and Deep Learning


## Training Process

But how can we check how well the training process is going?


- Plot the loss curve to see how the loss (error) develops over time (steps / epochs)




Ideal loss curve



Oscillating training  
- Bad training data  
- Learning rate too high



Disturbed training  
- Outliers in data  
- Missing values (NaN)




Overfitting on training data  
Bad generalizing for new data

Source: <https://developers.google.com/>  
Portions of this page are reproduced from work created and shared by Google and used according to terms described in the [Creative Commons 4.0 Attribution License](#).

# Neural Network Architectures

## Autoencoder


- Neural network for data compression (encoder) and reconstruction (decoder)
- Most important information (properties of features) get embedded in low-dimensional **latent space**
- Training aims to minimize the error during reconstruction
- Applications, e.g., dimensionality reduction or anomaly detection (large error = anomaly)



# Neural Network Architectures

## Recurrent Neural Network (RNN / LSTM / GRU)

- Neural network for modeling sequences
- Implements “memory” with the help of “feedback loops”
- Information no longer flows only from input to output (feed forward), but is also fed back into the neurons of previous hidden layers
- Variants of this enable, among other things, “long short-term memory” (LSTM)



$x$  – Input  
 $h$  – Hidden (Memory)  
 $o$  – Output  
 $U, V, W$  – according Weights

### Unfold...“roll up” over time:

At time  $t$ , the previous state  $t-1$  is incorporated via  $V$   
At time  $t+1$ , the previous state  $t$  is incorporated via  $V$

...


Source: fdeloche,  
<https://commons.wikimedia.org/w/index.php?curid=60109157>, CC BY-SA 4.0



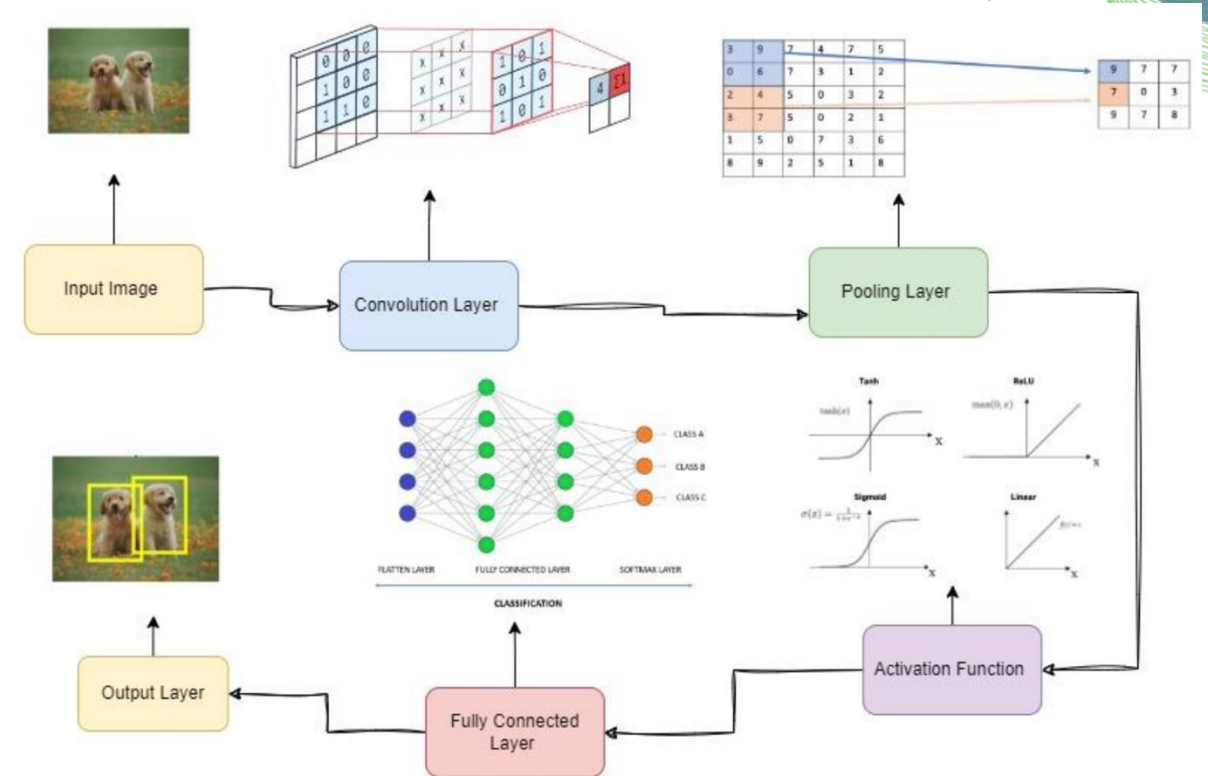
# Neural Network Architectures

## Convolutional Neural Network (CNN)

- Neural network with different layer types for extracting local to global features from input data
  - Convolutional Layer (Convolution)
  - Pooling Layer (Aggregation)
  - Dropout Layer (Regularization)
  - Fully Connected Layer (Classification)
- Application mostly in image processing (object recognition, classification)



Source: Cmglee,  
<https://commons.wikimedia.org/w/index.php?curid=104937230>, CC BY-SA 4.0



Source: Taye, M. M. (2023). Theoretical Understanding of Convolutional Neural Network: Concepts, Architectures, Applications, Future Directions. *Computation*, 11(3), 52. <https://doi.org/10.3390/computation11030052>, OpenAccess


# Neural Network Architectures

## Convolutional Neural Network (CNN)

- Convolutional Layer – not fully connected
- Kernel moves over input and extracts features by applying element-wise multiplications and sums  
→ kernels are small weight matrices that are learned during training

0	2	0
2	1	2
0	2	0

3x3 Kernel



- Pooling Layer aggregates input (e.g., max, avg, ...), introduces robustness to variations in the data
- Dropout randomly removes connections during training to prevent overfitting

Quellen:

Vincent Dumoulin and Francesco Visin (2018). A guide to convolution arithmetic for deep learning. *arXiv*. <https://doi.org/10.48550/arXiv.1603.07285>

Vincent Dumoulin, Francesco Visin, [https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic), MIT

3	15	1	13
9	7	0	10
11	5	5	3
1	8	9	6

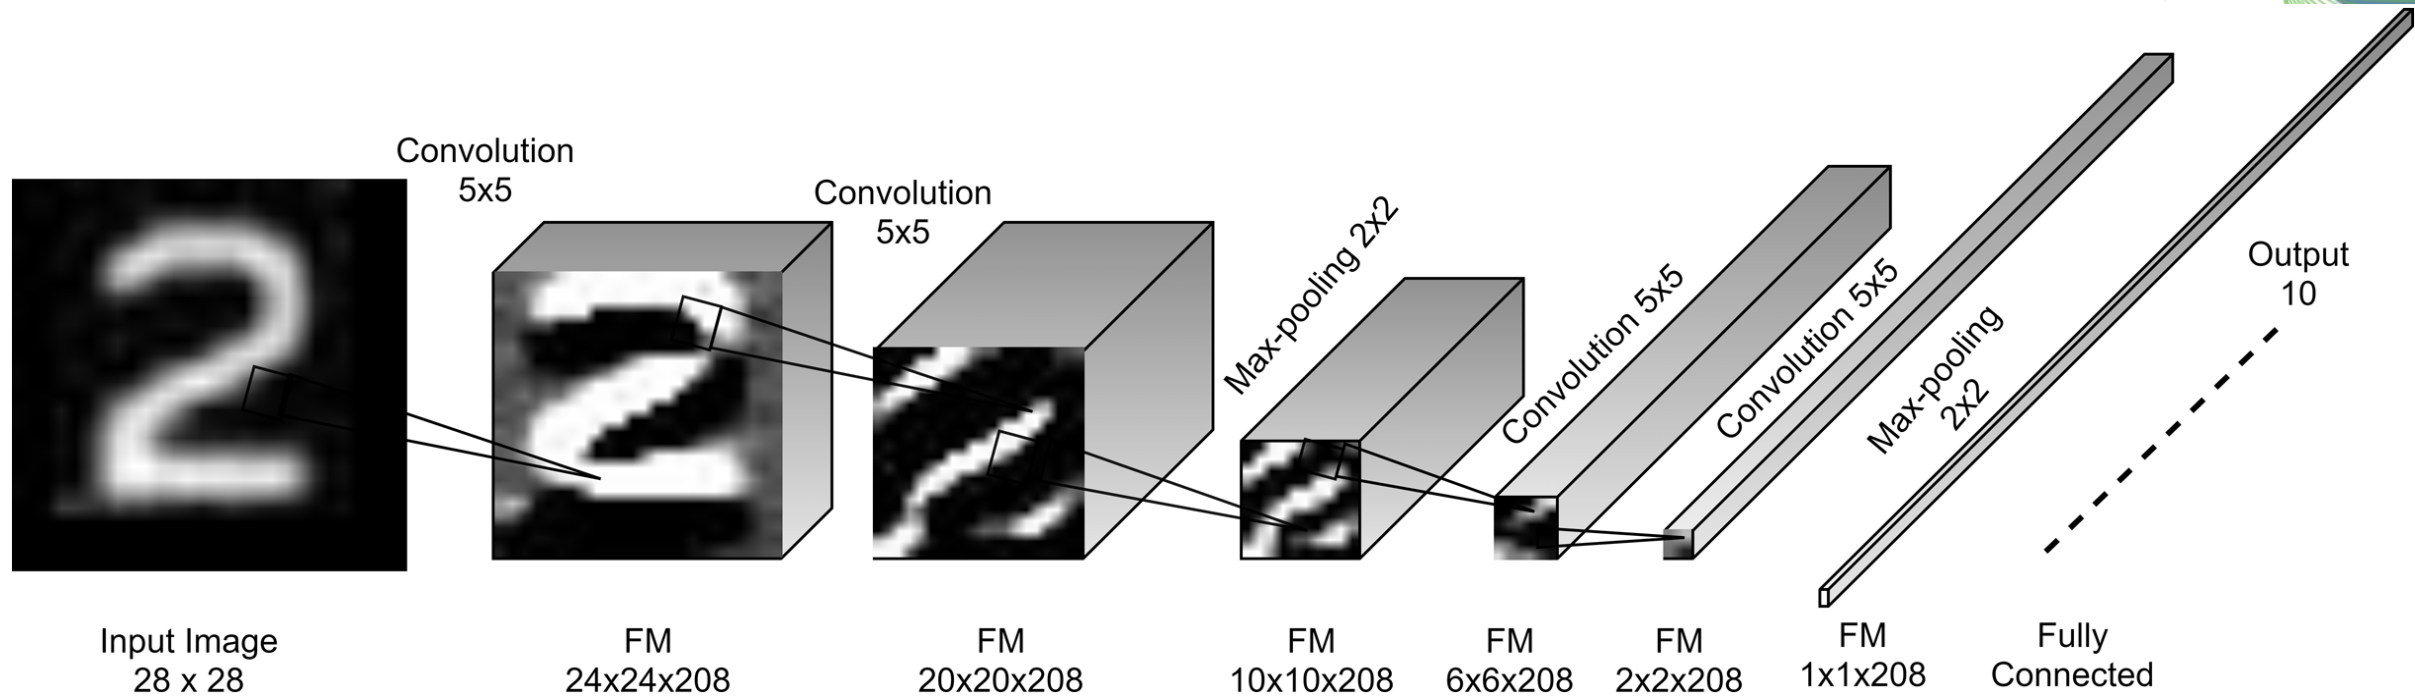
Max pooling

15	13
11	9

# Neural Network Architectures

## Convolutional Neural Network (CNN)

- Convolutional Neural Network – the full image




Source: Saleemdeen M, Ertürk S. 2021. Full depth CNN classifier for handwritten and license plate characters recognition. PeerJ Computer Science 7:e576  
<https://doi.org/10.7717/peerj-cs.576>, CC-BY4.0



# Neural Network Architectures

## Transformer

- Complex, multi-component, multi-layer deep learning-based model for generative AI, sequence modelling...
- Consists of various components, grouped into encoders and decoders




Quelle: Vaswani et al (2017), <https://arxiv.org/abs/1706.03762>

# Neural Network Architectures

## Transformer

### Transformer architecture explained





Source: Vaswani et al (2017), <https://arxiv.org/abs/1706.03762>

# Neural Network Architectures

## Transformer

Transformer – High level view on the process



# Practical exercise

- Build and train a neural network for a regression problem
- Covering model training, hyperparameter optimization, monitoring the training process, model evaluation
- Leveraging the Python libraries
  - [pytorch](#) – open-source deep learning framework
  - [optuna](#) – open-source framework for hyperparameter optimization