

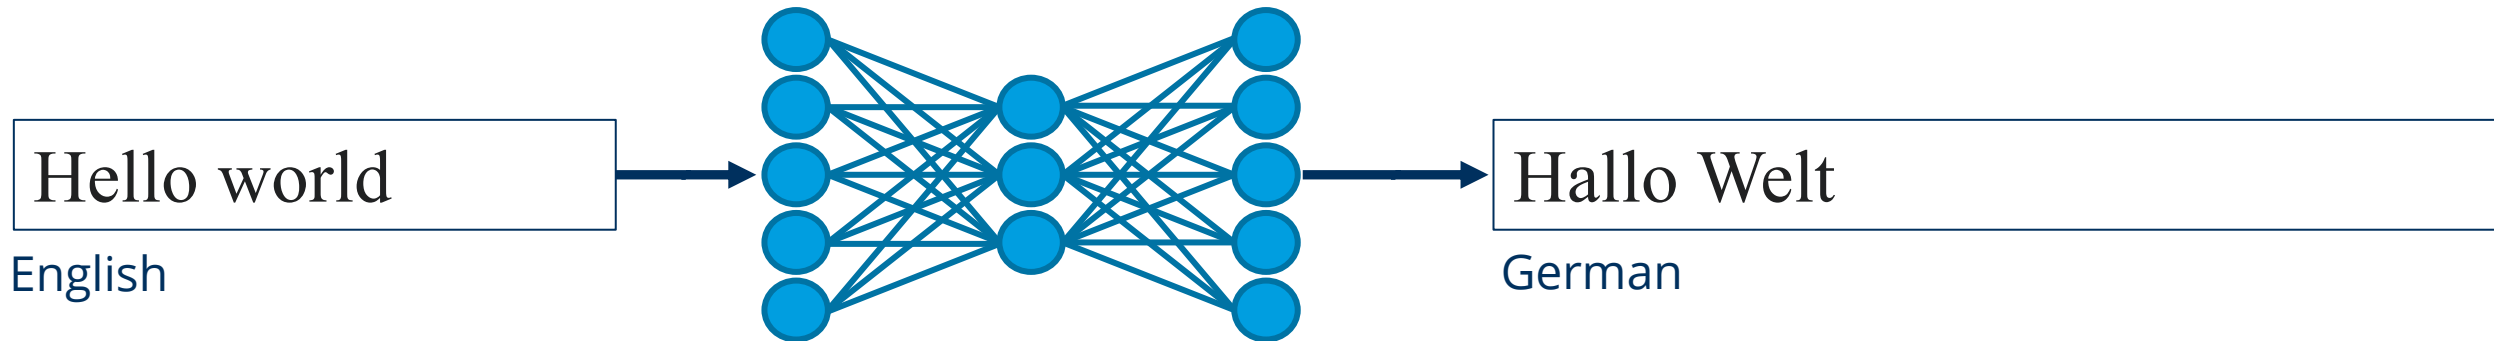
Using Large Language Models for Data Analysis

Robert Haase

These slides can be reused under the terms of the [CC-BY4.0](https://creativecommons.org/licenses/by/4.0/) license.

Large Language Models (LLMs)

Text-to-text, translation

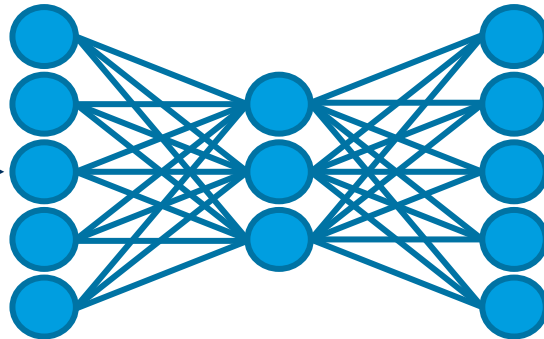


Large Language Models (LLMs)

Text-to-text, translation, code generation

Open hela-cells.tif

English



```
from skimage.io import imread  
image = imread("hela-cells.tif")
```

Python

Published as a conference paper at ICLR 2024

**SWE-BENCH: CAN LANGUAGE MODELS RESOLVE
REAL-WORLD GITHUB ISSUES?**

Carlos E. Jimenez^{* 1,2} John Yang^{* 1,2} Alexander Wettig^{1,2}
Shunyu Yao^{1,2} Kexin Pei³ Ofir Press^{1,2} Karthik Narasimhan^{1,2}

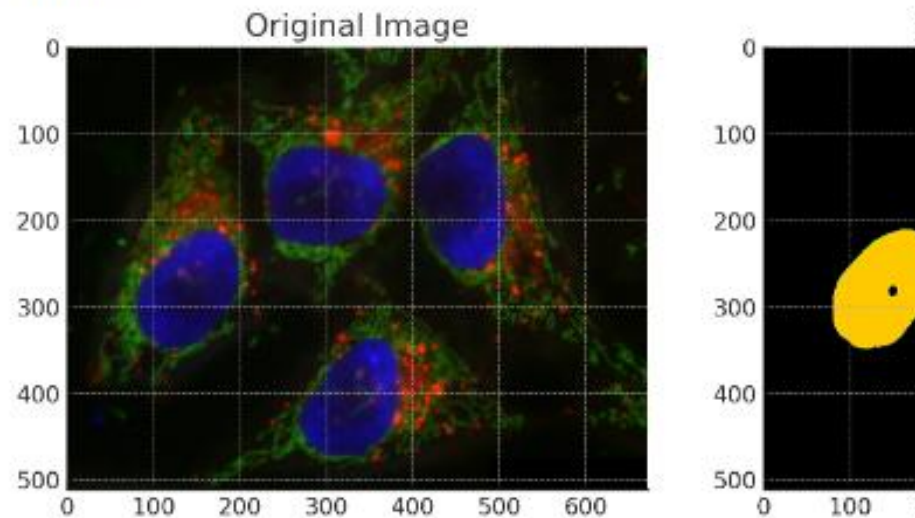
¹Princeton University ²Princeton Language and Intelligence ³University of Chicago

Bio-image Analysis using Large-L

Analysing images in the browser, e.g. using ChatC

Prompt:
“Please segment
the nuclei in the
blue channel of
the image”

ChatGPT



The segmentation of the nuclei has been performed on the blue channel
the figure shows the segmented nuclei in different colors, each corres
the field of view. Each color represents a unique label assigned to the
image. [~]



Analysis

```
from skimage import io, filters, measure, morphology
import matplotlib.pyplot as plt
import numpy as np

# Load the image
image = io.imread('/mnt/data/hela-cells-8bit.png')

# Since it's a color image, we'll just take the blue channel
blue_channel = image[:, :, 2]

# Threshold the image to get the binary image for nuclei
thresh = filters.threshold_otsu(blue_channel)
binary = blue_channel > thresh

# Remove small objects and perform morphological closing
cleaned = morphology.remove_small_objects(binary, 50)
closed = morphology.closing(cleaned, morphology.disk(3))

# Label the image
label_image = measure.label(closed)

# Plot the original image and the binary image
fig, ax = plt.subplots(1, 2, figsize=(12, 6))
ax[0].imshow(image)
ax[0].set_title('Original Image')
ax[1].imshow(label_image, cmap='nipy_spectral')
ax[1].set_title('Segmented Nuclei')
plt.show()
```

Reproducibility: Python

Technical: I can rerun the workflow because it is well documented.

Inter-personal: I understand a workflow and can explain all steps.

```
1 from skimage.io import imread, imsave
2 import stackview
3 import pyclesperanto as cle
4 import numpy as np
5
6 image = imread('data/lund.tif')
7
8 # Apply top hat filter to the image
9 processed_image = np.asarray(cle.top_hat_box(image, radius_x=10, radius_y=10, radius_z=0))
10 # Segment the image using voronoi otsu labeling
11 label_image = cle.voronoi_otsu_labeling(processed_image, spot_sigma=2, outline_sigma=2)
12
13 print(label_image.max())
```

Output: 285

Technical: 

Inter-personal: 

We cannot follow the workflow, because we do not see intermediate results.

Reproducibility: Python + Jupyter

Technical: I can rerun the workflow because it is well documented.

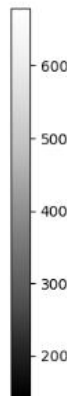
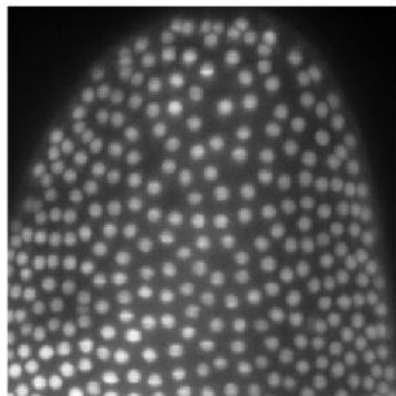
Inter-personal: I understand a workflow and can explain all steps.

```
[2]: image = imread('data/lund.tif')  
  
image.shape
```

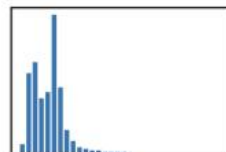
```
[2]: (100, 256, 256)
```

```
[3]: stackview.insight(image)
```

```
[3]:
```

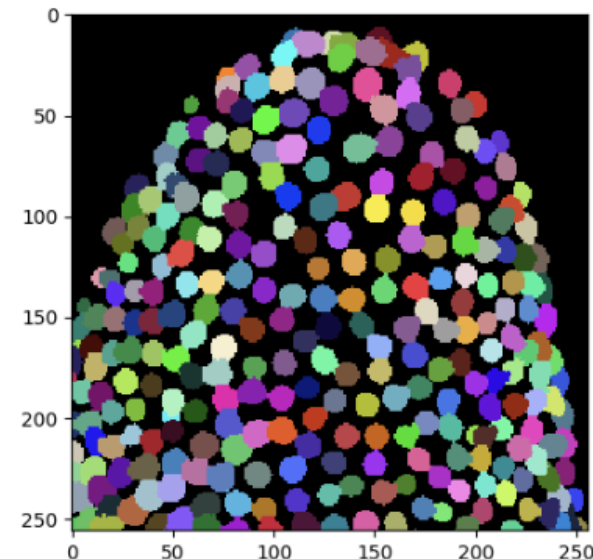


```
shape (100, 256, 256)  
dtype uint16  
size 12.5 MB  
min 125  
max 680
```



```
[12]: # Segment the image using voronoi otsu labeling  
label_image = cle.voronoi_otsu_labeling(processed_image, spot_sigma=2, outline_sigma=2)  
label_image
```

```
[12]:
```



cle_image

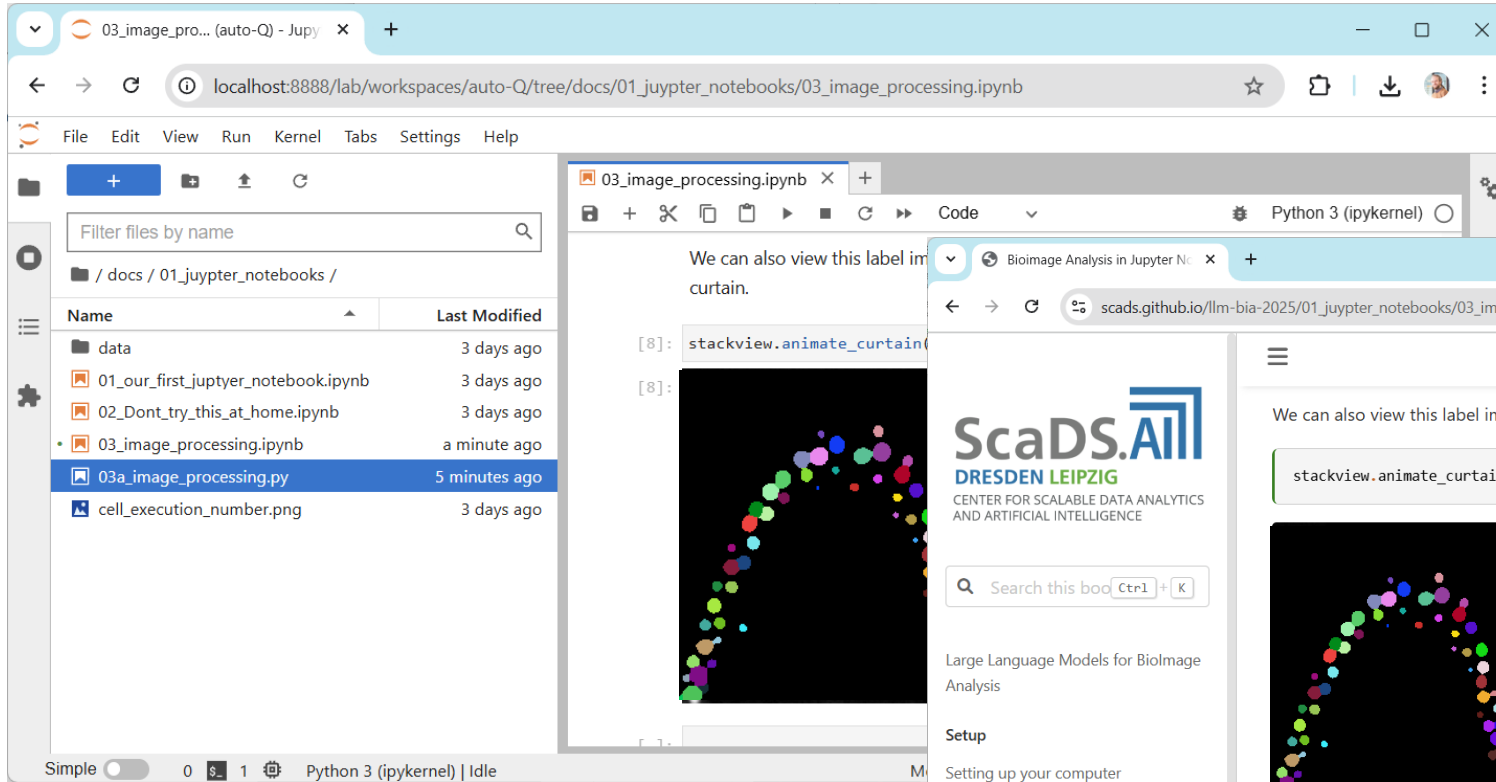
```
shape (100, 256, 256)  
dtype uint32  
size 25.0 MB  
min 0.0  
max 285.0
```

Technical: ☒

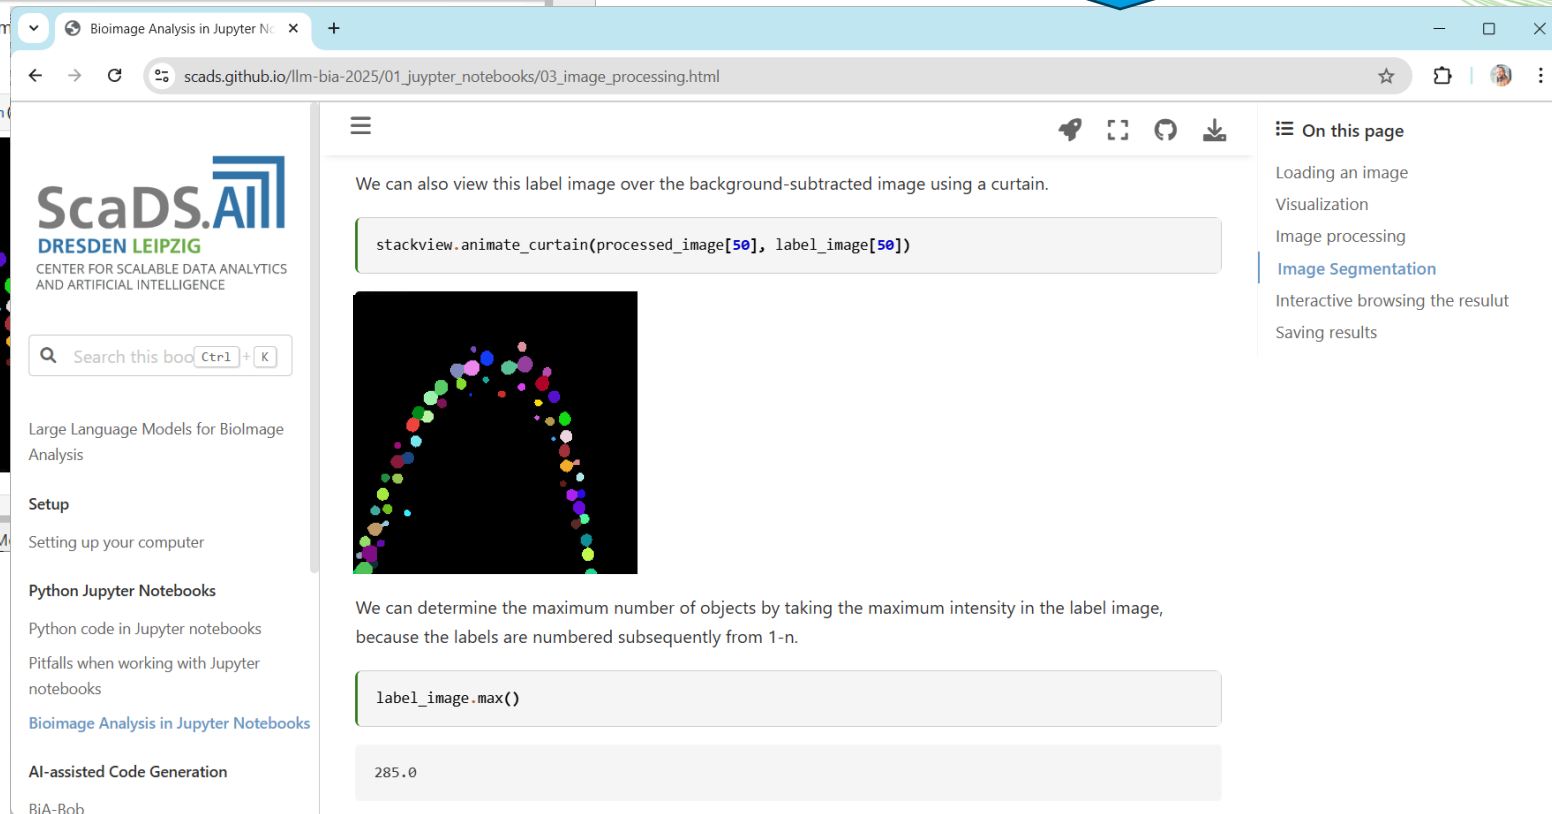
Inter-personal: ☒

Reproducibility : Python + Jupyter

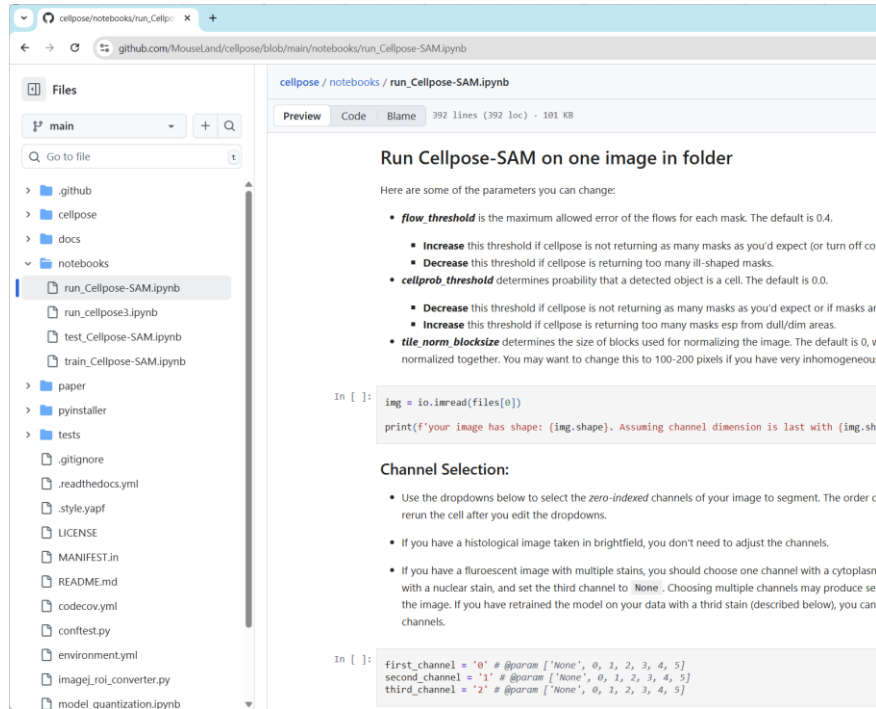
Other's view when
you share the
workflow online



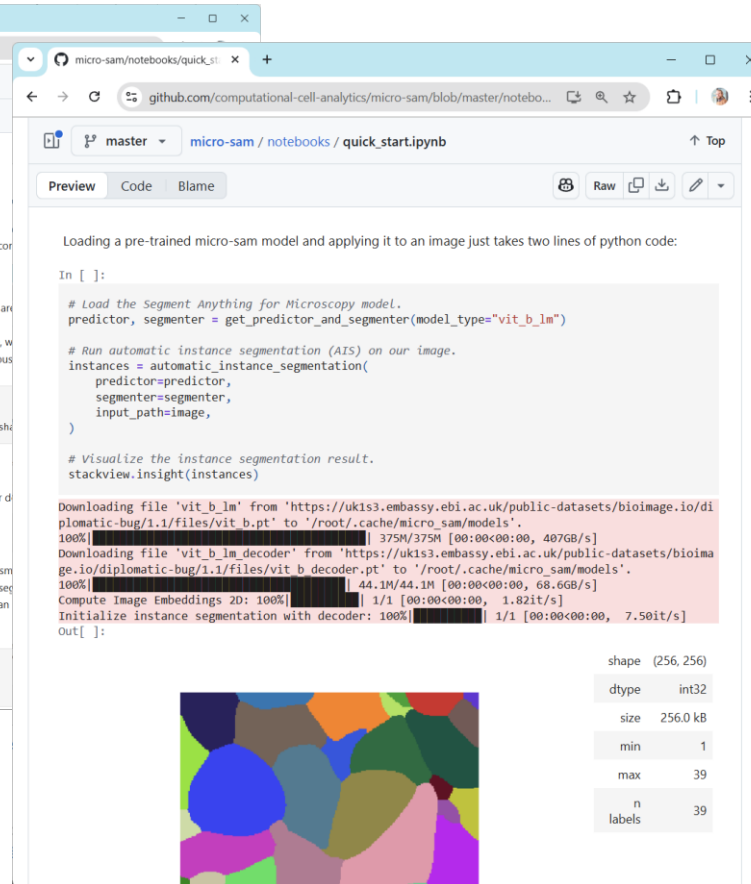
View on your laptop



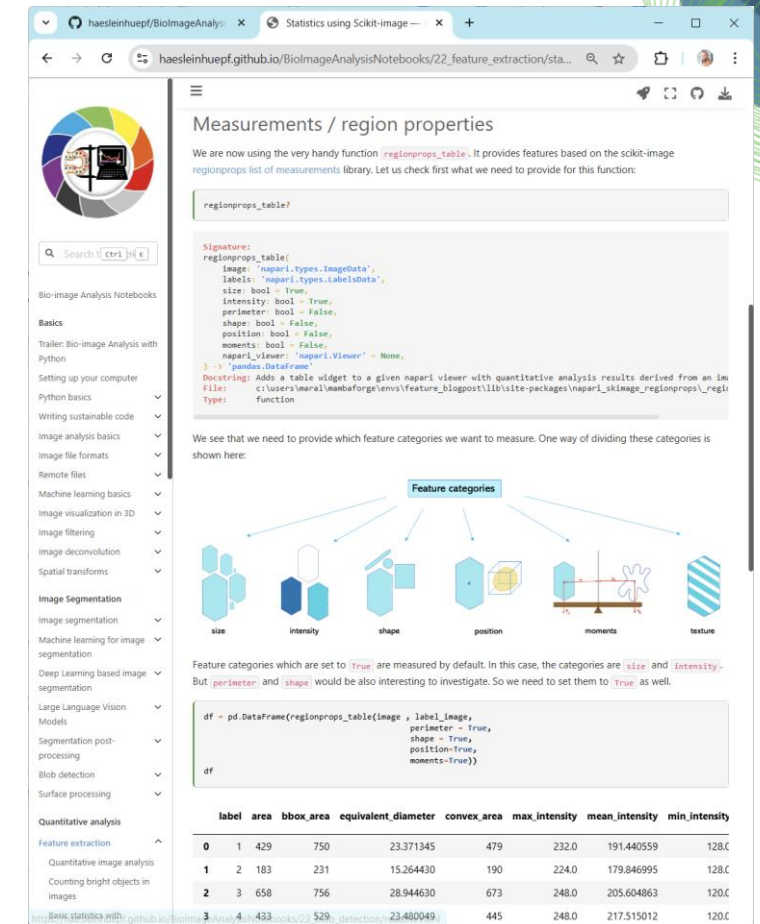
Reproducibility : Python + Jupyter



https://github.com/MouseLand/cellpose/blob/main/notebooks/run_Cellpose-SAM.ipynb

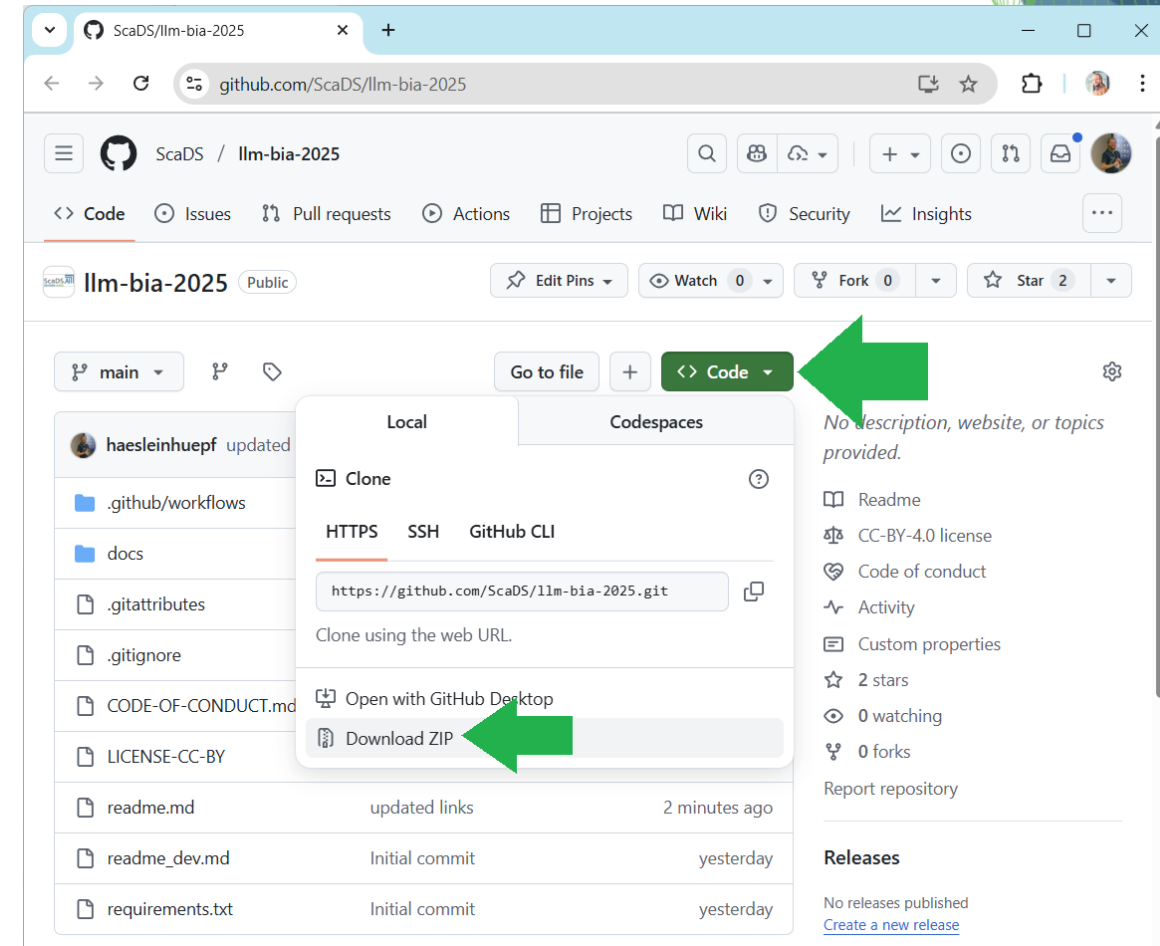
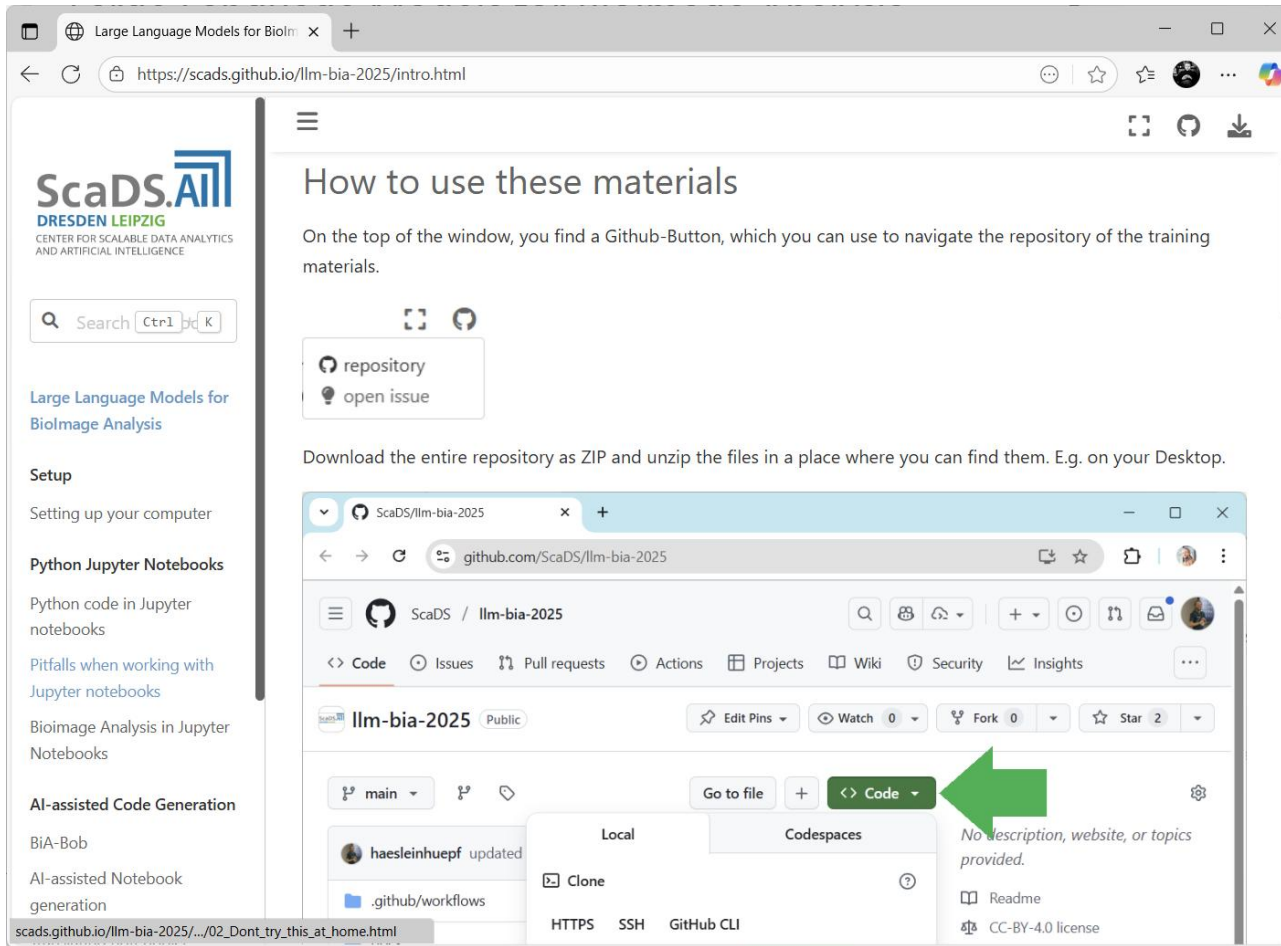


https://github.com/computational-cell-analytics/micro-sam/blob/master/notebooks/quick_start.ipynb



https://haesleinhuepf.github.io/BiolImageAnalysisNotebooks/22_feature_extraction/statistics_with_scikit_image.html

Training materials for today



Why code generation using LLMs?

Output of LLMs is often criticized for

- Limited reproducibility
- Hard to evaluate (automatically)

Output of LLM-generated code is

- Reproducible
- Can be executed and evaluated (automatically)

Say a

Say a

Write Python code that says "a".

A.

< 2/2 >

A!

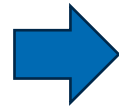
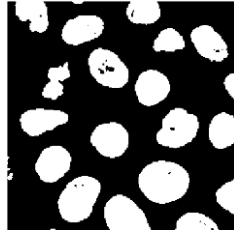
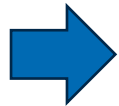
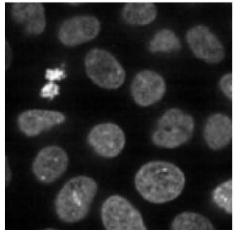
Certainly! Here's a simple Python code snippet that prints the letter "a":

```
python  
  
print("a")
```

No matter how often you execute this code, the output is always guaranteed to be „a“.

How good are LLMs for Bio-image Analysis?

Test case: segment the image and measure the average area of objects.



area	
0	955.0
1	31.0
2	815.0
3	1166.0
4	1135.0
⋮	



858.04

How good are LLMs for Bio-image Analysis?

Example test-case inspired by HumaEval (Chen et al 2021)

```
[1]: def workflow_segmentation_measurement_summary(image):  
    """  
    This function implements a workflow consisting of these steps:  
    * threshold intensity input image using Otsu's method  
    * label connected components  
    * measure area of the labeled objects  
    * determine mean area of all objects  
    """  
  
    import skimage  
    import numpy as np  
    binary_image = image > skimage.filters.threshold_otsu(image)  
    label_image = skimage.measure.label(binary_image)  
    stats = skimage.measure.regionprops(label_image)  
    areas = [s.area for s in stats]  
    return np.mean(areas)
```

Prompt

Reference
solution

```
[2]: def check(candidate):  
    import numpy as np  
  
    assert candidate(np.asarray([  
        [0,0,0,0,0],  
        [1,1,1,0,0],  
        [1,1,1,0,0],  
        [1,1,0,0,0],  
        [0,0,0,0,0],  
    ])) == 8
```

Unit test
(excerpt)

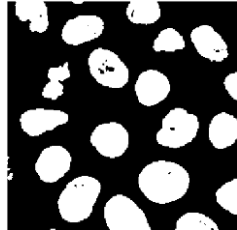
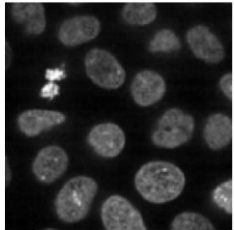
Open for your (!)
contributions

18 more currently on
their way...

We formulated 57 of
such test-cases (yet)

How good are LLMs for Bio-image Analysis?

Test case: segment the image and measure the average area of objects.



	area
0	955.0
1	31.0
2	815.0
3	1166.0
4	1135.0
...	



858.04

workflow_segmentation_measurement_summary



1.0	1.0	1.0	0.9	1.0	0.8	0.9	0.0	0.3	0.4	0.5	0.0	0.6	0.1	0.4	0.4	0.1	0.5	0.0	0.1	0.1
reference	claude-3-5-sonnet-20240620	gpt-4o-2024-05-13	gpt-4-turbo-2024-04-09	claude-3-opus-20240229	gpt-4-1106-preview	gemini-1.5-pro-001	gpt-4o-mini-2024-07-18	llama3-70b-instruct-q8_0	llama3-70b-instruct-q4_0	gpt-3.5-turbo-1106	gemini-1.5-flash-001	codegemma-7b-instruct-fp16	mixtral-8x22b-instruct-v0.1-q4_0	mixtral-8x7b-instruct-v0.1-q5_0	phi3-3.8b-mini-instruct-4k-fp16	codellama-70b-instruct-q4_0	gemini-pro	llama3-8b-instruct-fp16	command-r-plus-104b-q4_0	codellama



Unit-test pass-rate (n=10):


Large language models


Benchmarking LLMs for Bio-image Analysis

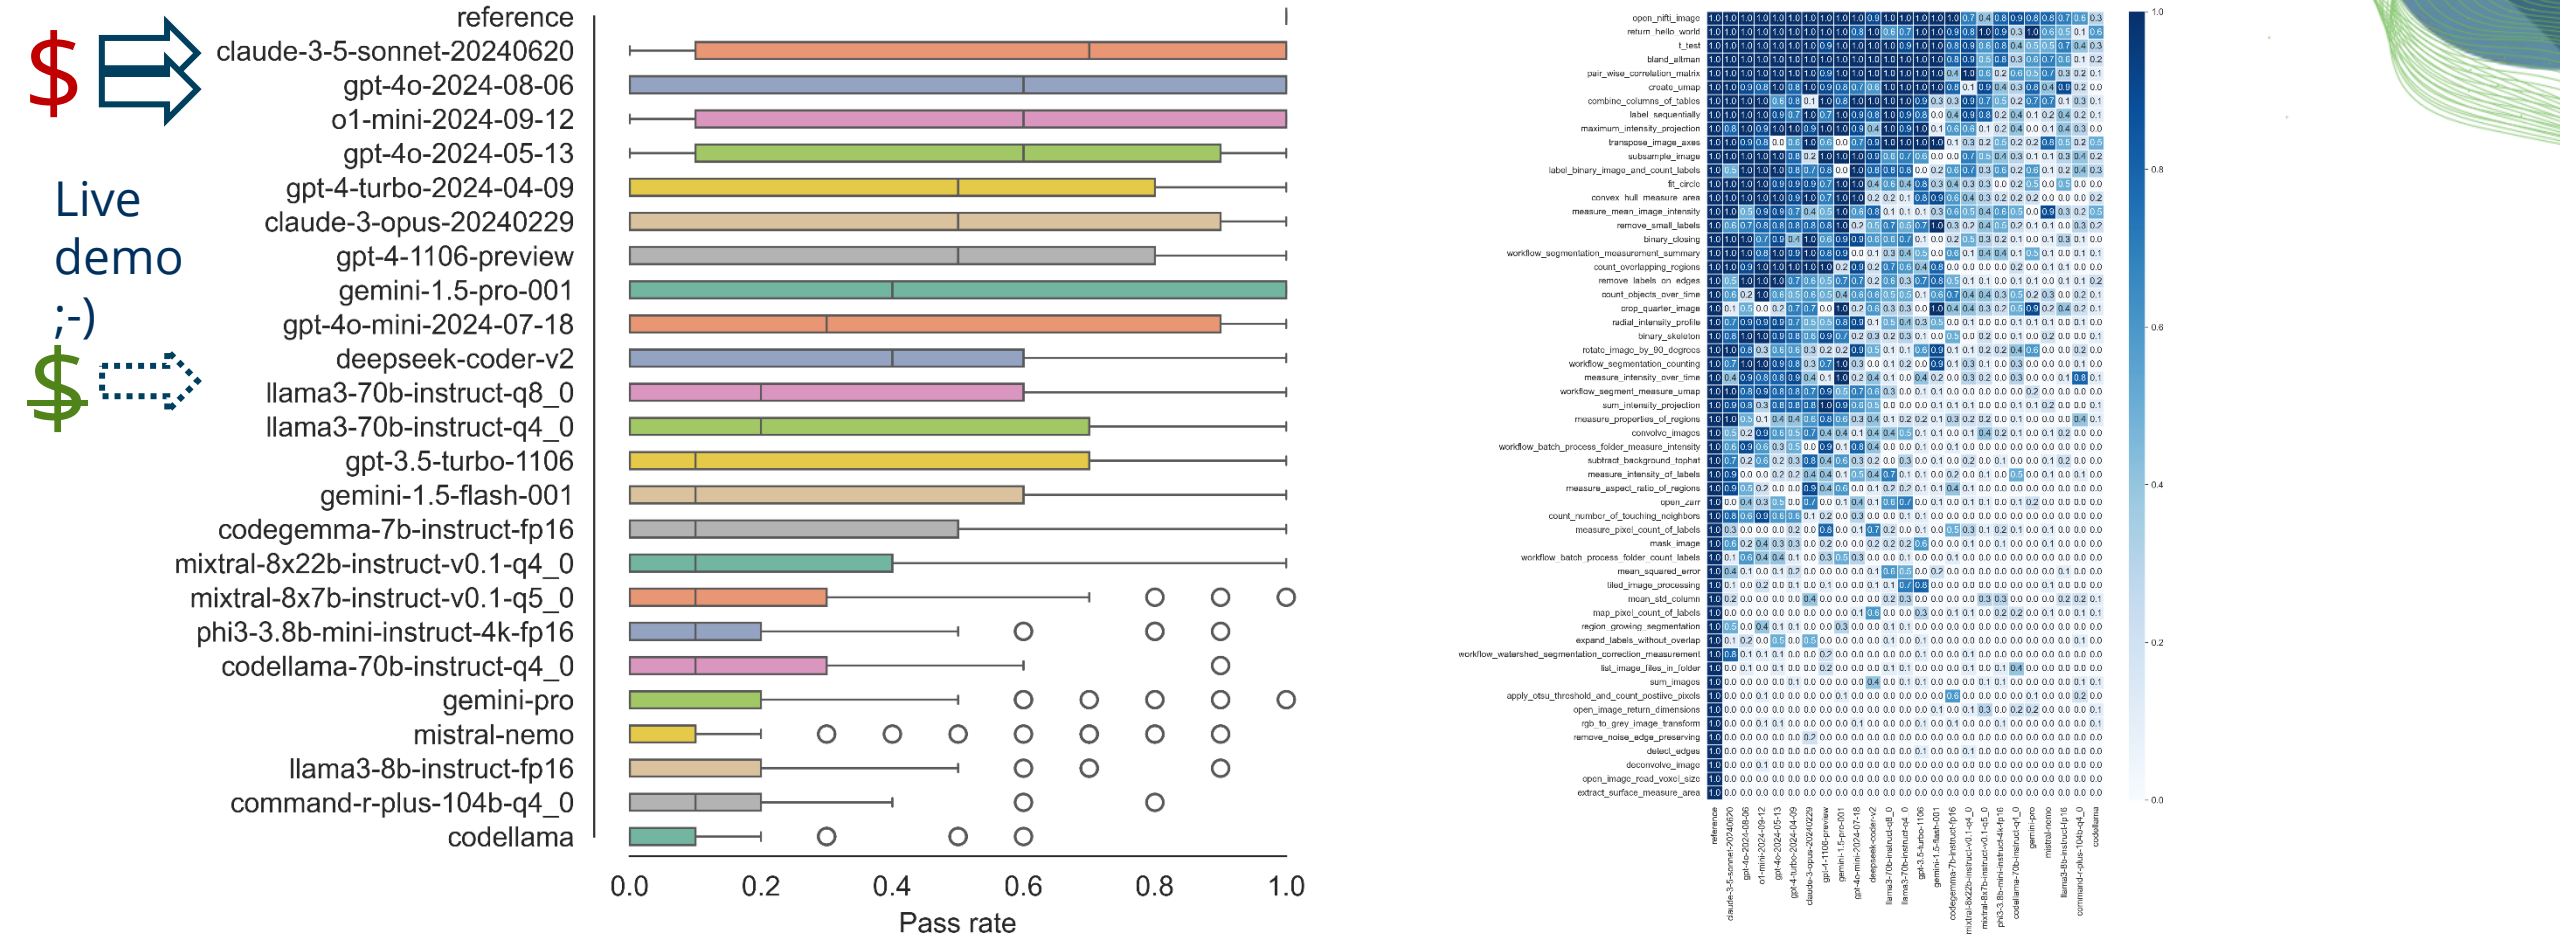
Summary: 57 use-cases (yet), 24 LLMs (yet), n=10

	reference
claude-3-5-sonnet-20240620	
gpt-4o-2024-08-06	
o1-mini-2024-09-12	



 reference
claude-3-5-sonnet-20240620
gpt-4o-2024-08-06
o1-mini-2024-09-12
gpt-4o-2024-05-13
gpt-4-turbo-2024-04-09
claude-3-opus-20240229
gpt-4-1106-preview
gemini-1.5-pro-001
gpt-4o-mini-2024-07-18
deepseek-coder-v2
llama3-70b-instruct-q8_0
llama3-70b-instruct-q4_0
gpt-3.5-turbo-1106
gemini-1.5-flash-001
codegemma-7b-instruct-fp16
mixtral-8x22b-instruct-v0.1-q4_0
mixtral-8x7b-instruct-v0.1-q5_0
phi3-3.8b-mini-instruct-4k-fp16
codellama-70b-instruct-q4_0
gemini-pro
mistral-nemo
llama3-8b-instruct-fp16
command-r-plus-104b-q4_0
codellama



Prompt Engineering

To further improve
LLMs capabilities

Prompt Engineering

Modify a prompt until it works.

Write python code to open the image file 'data/blobs.tif'
and threshold it using Otsu's method.
Apply connected component labeling to the result to
retrieve a label image. Visualize the label image.



Here is the python code to achieve the task:

```
import cv2
import numpy as np
import matplotlib.pyplot as plt

# Load the image
img = cv2.imread('data/blobs.tif', 0)

# Apply Otsu's thresholding
_, thresholded = cv2.threshold(img, 0, 255, cv2.THRESH_BINARY + cv2.THRESH_OTSU)

# Apply connected component labeling
num_labels, labels = cv2.connectedComponents(thresholded)

# Visualize the label image
plt.figure(figsize=(10, 10))
plt.imshow(labels, cmap='nipy_spectral')
plt.colorbar()
plt.show()
```

Make sure you have the necessary libraries (OpenCV, NumPy, and Matplotlib) installed in your Python environment before running the code.

Prompt Engineering

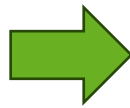
Modify a prompt until it works.

Write python code to open the image file 'data/blobs.tif' and threshold it using Otsu's method.

Apply connected component labeling to the result to retrieve a label image. Visualize the label image.

Do not explain the solution.

Only write Python code and short comments are ok.



```
import cv2
import numpy as np

# Read the image
img = cv2.imread('data/blobs.tif', cv2.IMREAD_GRAYSCALE)

# Apply Otsu's thresholding
_, thresh = cv2.threshold(img, 0, 255, cv2.THRESH_BINARY + cv2.THRESH_OTSU)

# Apply connected component labeling
num_labels, label_image = cv2.connectedComponents(thresh)

# Visualize the label image
cv2.imshow('Label Image', label_image.astype(np.uint8) * 50)
cv2.waitKey(0)
cv2.destroyAllWindows()
```

Prompt Engineering

Modify a prompt until it works.

„Long“ Context

Write python code to open the image file 'data/blobs.tif' and threshold it using Otsu's method.

Apply connected component labeling to the result to retrieve a label image. Visualize the label image.

Do not explain the solution.

Only write Python code and short comments are ok.

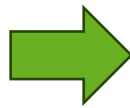
Assume the code is executed in a Jupyter notebook.

Do not use any of these Python libraries:

* OpenCV

* PIL

* Pillow



```
import matplotlib.pyplot as plt
from skimage import io, filters, measure

# Load the image
image = io.imread('data/blobs.tif')

# Threshold using Otsu's method
thresh = filters.threshold_otsu(image)
binary_image = image > thresh

# Connected component labeling
label_image = measure.label(binary_image)

# Visualize the label image
plt.figure(figsize=(8,8))
plt.imshow(label_image, cmap='nipy_spectral')
plt.axis('off')
plt.show()
```


Scientific code writing: bia-bob



- First steps: Keep it simple

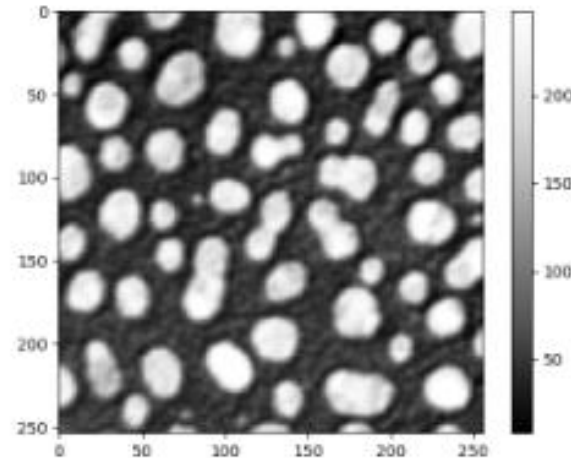
```
%bob Load the blobs.tif image file and show it.
```

I will load the `blobs.tif` image file and then display it using the appropriate visualization function.

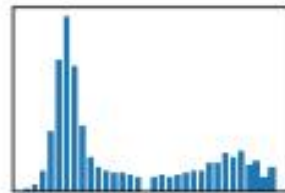
```
from skimage.io import imread
import stackview

# Load the image file
filename = 'blobs.tif'
image = imread(filename)

# Display the image
stackview.insight(image)
```



shape (254, 256)
dtype uint8
size 63.5 kB
min 8
max 248



Scientific code writing: bia-bob

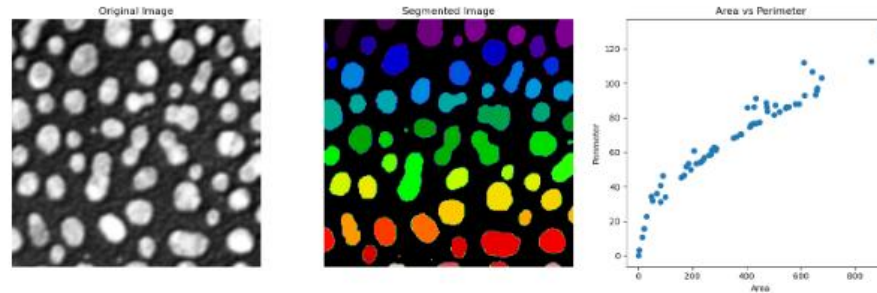


- After learning *how to prompt* for code: try more complex workflows

```
[1]: import bia_bob
```

```
[2]: %%bob
* Load blobs.tif from disk
* Segment the bright objects on the dark background
* Measure area and perimeter of the objects
* Show the original image, the segmented image and an
  area versus perimeter plot side-by-side using subfigures.
```

This notebook may contain text, code and images generated by artificial intelligence. Used model: claude-3-5-sonnet-20240620, vision model: claude-3-5-sonnet-20240620, endpoint: None, bia-bob version: 0.21.0.. Do not enter sensitive or private information and verify generated contents according to good scientific practice. Read more: <https://github.com/haesleinhuepf/bia-bob#disclaimer>



	label	area	perimeter
0	1	433.0	91.254834
1	2	185.0	53.556349
2	3	658.0	95.698485
3	4	434.0	76.870058
4	5	477.0	83.798990

```
import numpy as np
from skimage import io, filters, measure, segmentation
import matplotlib.pyplot as plt
import pandas as pd
import pyclesperanto_prototype as cle
import stackview

# Load the image
image = io.imread('blobs.tif')

# Segment the bright objects
thresh = filters.threshold_otsu(image)
binary = image > thresh

# Label the objects
labels = measure.label(binary)

# Measure properties
properties = ['label', 'area', 'perimeter']
df = pd.DataFrame(measure.regionprops_table(labels, properties=properties))

# Create the figure with subfigures
fig, (ax1, ax2, ax3) = plt.subplots(1, 3, figsize=(15, 5))

# Original image
ax1.imshow(image, cmap='gray')
ax1.set_title('Original Image')
ax1.axis('off')

# Segmented image
ax2.imshow(labels, cmap='nipy_spectral')
ax2.set_title('Segmented Image')
ax2.axis('off')

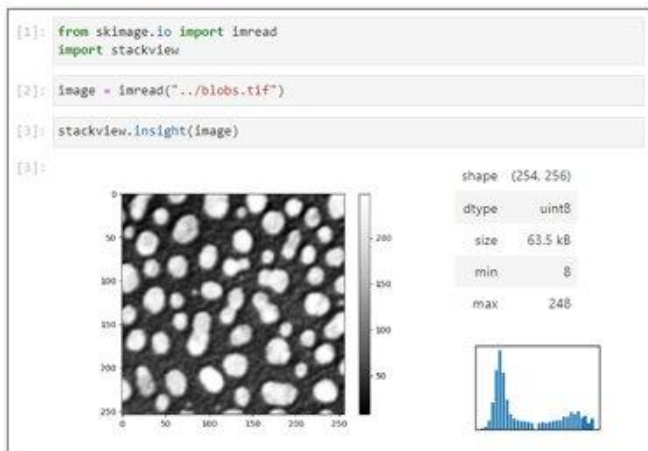
# Scatter plot
ax3.scatter(df['area'], df['perimeter'])
ax3.set_xlabel('Area')
ax3.set_ylabel('Perimeter')
ax3.set_title('Area vs Perimeter')

plt.tight_layout()
plt.show()

# Display the dataframe
print(df.head())
```

Optimizing documentation

E.g. to make code more readable



```
from bia_bob import bob
```

`%bob` modify the notebook `mod_notebook.ipynb` by adding explanatory markdown cells between the code cells. Also add a markdown cell at the very beginning that introduces the entire notebook shortly.

The modified notebook has been saved as `mod_notebook_1.ipynb`.

Introduction

In this notebook, we will demonstrate how to load a bioimage, visualize it using stackview, and explore the image interactively. We will be using an image of blobs for this example. The steps include loading the image and visualizing it using different stackview functionalities.

Step 1: Import necessary libraries

We will start by importing the necessary libraries. In this case, `skimage.io` for loading the image and `stackview` for visualization.

```
from skimage.io import imread
import stackview
```

Step 2: Load the image

Next, we will load the image file into a variable. The image used in this notebook is `blobs.tif`.

```
image = imread("../blobs.tif")
```

Step 3: Visualize the image using stackview insight

We will use `stackview.insight` to visualize the loaded image. This function allows for inspecting the image.

```
stackview.insight(image)
```

shape (254, 256)
dtype uint8
size 63.5 kB
min 8
max 248

Translating notebooks



"Please translate
notebook
<xyz.ipynb> to
<language>."



10_image_data.ipynb

Multidimensional image stacks

Multidimensional image data can be handled in a similar multi-channel image data.

3-dimensional image stacks

There are also images with three spatial dimensions: X, Y, and Z. Typical examples in microscopy and in medical imaging: a look at an Magnetic Resonance Imaging (MRI) data

```
from skimage.io import imread
from stackview import imshow
import matplotlib.pyplot as plt
image_stack = imread('.../data/Haase_MRT_tf13d1.tif')
```

Image slicing

Since three dimensions are spatial dimensions, we can also make orthogonal to the image plane and corresponding to the anatomical planes. To orient the images correctly, we can transpose their axes by adding '.T' by the end.

```
sagittal = image_stack[:, :, 128].T
coronal = image_stack[:, 128, :].T
transverse = image_stack[96]
```

```
fig, axs = plt.subplots(1, 3, figsize=(15,15))

# show orthogonal planes
axs[0].imshow(sagittal, cmap='Greys_r')
axs[0].set_title('Sagittal')

axs[1].imshow(coronal, cmap='Greys_r')
axs[1].set_title('Coronal')

axs[2].imshow(transverse, cmap='Greys_r')
axs[2].set_title('Transverse');
```



Videos

If an image dataset has a temporal dimension, we call it a video. Processing videos works similar to multi-channel images and image stacks. Let's open a microscopy dataset showing yeast



10_image_data.de.ipynb

Multidimensionale Bildstapel

Multidimensionale Bilddaten können ähnlich wie mehrkanalige Bilddaten behandelt werden.

Dreidimensionale Bildstapel

Es gibt auch Bilder mit drei räumlichen Dimensionen: X, Y und Z. Typische Beispiele finden sich in der Mikroskopie und in der medizinischen Bildgebung. Schauen wir uns ein Magnetresonanztomographie (MRT) Datensatz an:

```
[1]: from skimage.io import imread
      from stackview import imshow
      import matplotlib.pyplot as plt
      image_stack = imread('.../data/Haase_MRT_tf13d1.tif')
```

Bildschnitte

Da alle drei Dimensionen räumliche Dimensionen sind, können wir auch Schnitte orthogonal zur Bildebene machen, die den Anatomischen Ebenen entsprechen. Um die Bilder korrekt zu orientieren, können wir ihre Achsen transponieren, indem wir '.T' hinzufügen.

```
[2]: sagittal = image_stack[:, :, 128].T
      coronal = image_stack[:, 128, :].T
      transverse = image_stack[96]
```

```
fig, axs = plt.subplots(1, 3, figsize=(15,15))

# orthogonale Ebenen anzeigen
axs[0].imshow(sagittal, cmap='Greys_r')
axs[0].set_title('Sagittal')

axs[1].imshow(coronal, cmap='Greys_r')
axs[1].set_title('Coronal')

axs[2].imshow(transverse, cmap='Greys_r')
axs[2].set_title('Transversal');
```



Videos

Wenn ein Bilddatensatz eine zeitliche Dimension hat, nennen wir ihn ein Video. Die Verarbeitung von Videos funktioniert ähnlich wie die mehrkanaliger Bilder und Bildstapel. Öffnen wir einen Mikroskopie-



10_image_data.es.ipynb

Pilas de imágenes multidimensionales

Los datos de imágenes multidimensionales se pueden manejar de manera similar a los datos de imágenes multicanal.

Pilas de imágenes tridimensionales

También hay imágenes con tres dimensiones espaciales: X, Y y Z. Puedes encontrar ejemplos típicos en microscopía y en imágenes médicas. Echemos un vistazo a un conjunto de datos de imágenes por Resonancia Magnética (IRM):

```
[1]: from skimage.io import imread
      from stackview import imshow
      import matplotlib.pyplot as plt
      image_stack = imread('.../data/Haase_MRT_tf13d1.tif')
```

Corte de imágenes

Como las tres dimensiones son dimensiones espaciales, también podemos hacer cortes ortogonales al plano de la imagen y correspondientes a Planos Anatómicos. Para orientar correctamente las imágenes, podemos transponer sus ejes añadiendo '.T' al final.

```
[2]: sagittal = image_stack[:, :, 128].T
      coronal = image_stack[:, 128, :].T
      transverse = image_stack[96]
```

```
fig, axs = plt.subplots(1, 3, figsize=(15,15))

# mostrar planos ortogonales
axs[0].imshow(sagittal, cmap='Greys_r')
axs[0].set_title('Sagittal')

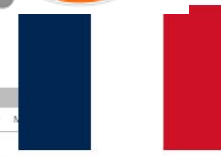
axs[1].imshow(coronal, cmap='Greys_r')
axs[1].set_title('Coronal')

axs[2].imshow(transverse, cmap='Greys_r')
axs[2].set_title('Transversal');
```



Videos

Si un conjunto de datos de imágenes tiene una dimensión temporal, lo llamamos video. Procesar videos funciona de manera similar a las imágenes multicanal u a las pilas de imágenes. Vamos a abrir un



10_image_data.fr.ipynb

Piles d'images multidimensionnelles

Les données d'images multidimensionnelles peuvent être traitées de manière similaire aux données d'images multi-canaux.

Piles d'images tridimensionnelles

Il existe aussi des images avec trois dimensions spatiales: X, Y et Z. Vous trouvez des exemples typiques en microscopie et en imagerie médicale. Regardons un ensemble de données d'imagerie par résonance magnétique (IRM):

```
[1]: from skimage.io import imread
      from stackview import imshow
      import matplotlib.pyplot as plt
      image_stack = imread('.../data/Haase_MRT_tf13d1.tif')
```

Coupe d'image

Comme les trois dimensions sont des dimensions spatiales, nous pouvons également faire des coupes orthogonales au plan de l'image et correspondant aux plans anatomiques. Pour orienter correctement les images, nous pouvons transposer leurs axes en ajoutant '.T' à la fin.


```
[2]: sagittal = image_stack[:, :, 128].T
      coronal = image_stack[:, 128, :].T
      transverse = image_stack[96]
```

```
fig, axs = plt.subplots(1, 3, figsize=(15,15))

# montrer des plans orthogonaux
axs[0].imshow(sagittal, cmap='Greys_r')
axs[0].set_title('Sagittal')

axs[1].imshow(coronal, cmap='Greys_r')
axs[1].set_title('Coronal')

axs[2].imshow(transverse, cmap='Greys_r')
axs[2].set_title('Transverse');
```



Vidéos

Si un ensemble de données d'image a une dimension temporelle, nous l'appelons une vidéo. Le traitement des vidéos fonctionne de

GPU-Accelerating code



Original:

```
[7]: %%acc
from skimage.io import imread
from skimage.filters import threshold_otsu
from skimage.measure import label
from skimage.color import label2rgb
import stackview

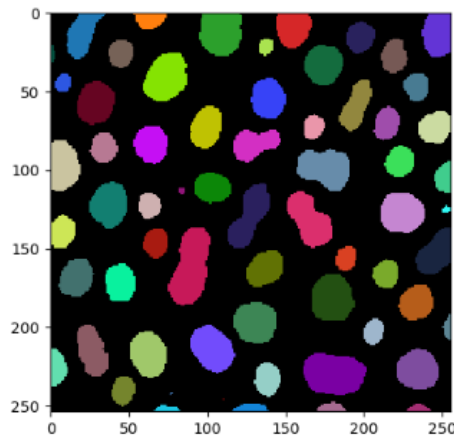
# Load the image
image = imread('blobs.tif')

# Apply global Otsu threshold to segment the image
binary_image = image > threshold_otsu(image)

# Label the objects
labels = label(binary_image)

# Display the color labels
stackview.insight(labels)
```

[7]:



shape	(254, 256)
dtype	int32
size	254.0 kB
min	0
max	64

Replacement:

```
[8]: import pyclesperanto_prototype as cle
import stackview

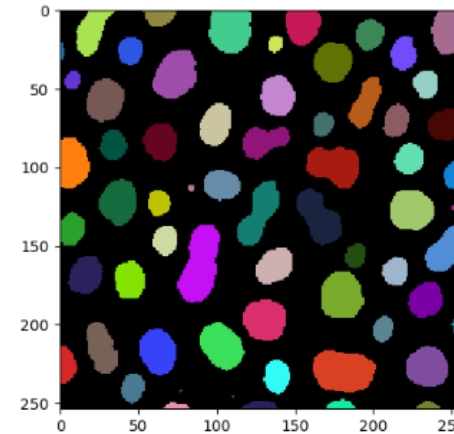
# Load the image
image = cle.imread('blobs.tif')

# Apply global Otsu threshold to segment the image
binary_image = cle.threshold_otsu(image)

# Label the objects
labels = cle.connected_components_labeling_box(binary_image)

# Display the color labels
stackview.insight(labels)
```

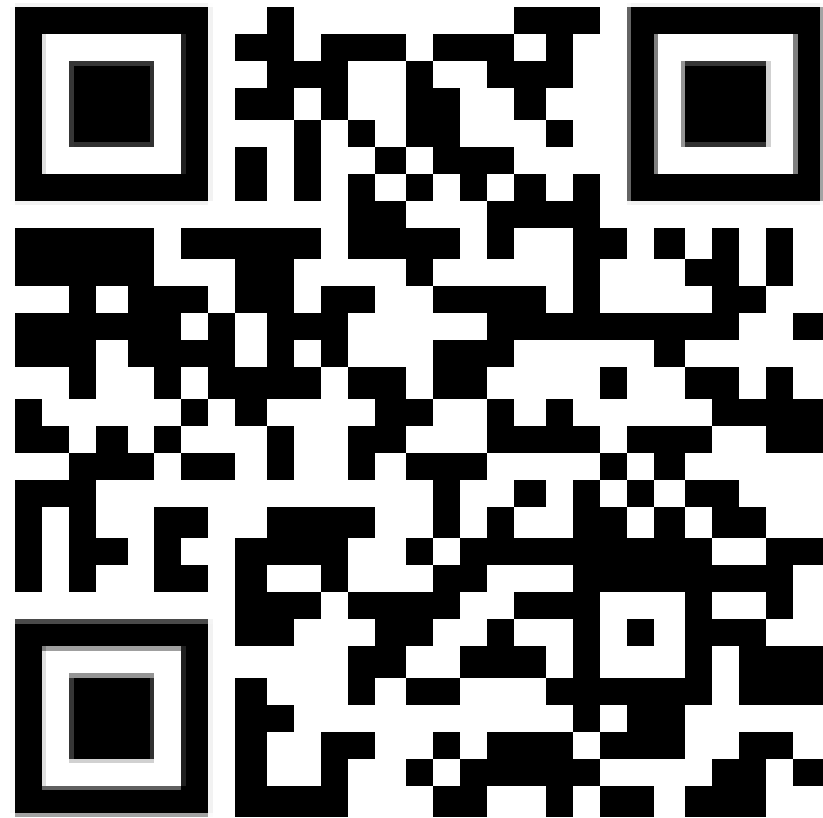
[8]:



It's our
responsibility
that results are
comparable!

https://link.springer.com/chapter/10.1007/978-3-030-76394-7_5

Exercises



<https://scads.github.io/llm-bia-2025/>

Bia-Bob: code generation

The image shows two overlapping browser windows. The left window displays the Bia-Bob demo website at scads.github.io/EMBOBioImage2025/02_bia_bob/bia-bob-demo.html. The right window shows a JupyterLab interface at localhost:8888/lab/tree/docs/02_bia_bob/bia-bob-demo.ipynb.

Website Content:

- Setup**: Setting up your computer
- Python Jupyter Notebooks**: Python code in Jupyter notebooks, Pitfalls when working with Jupyter notebooks, Bioimage Analysis in Jupyter Notebooks
- AI-assisted Code Generation**: **BiA-Bob**, AI-assisted Notebook generation, Translating notebooks, Statistical testing, Under the hood: Configure bia-bob's behaviour through system messages
- Links**: EMBO Practical Course Advanced

BiA-Bob Description:

BiA-Bob is a Python library that connects Jupyter with large language analysis instructions. Commonly, it uses OpenAI's/GPT models LLM infrastructure of [ScaDS.AI](#) in the [National Compute Center](#).

When using bia-bob to prompt LLMs, the prompts you enter in the LLM provider. Do not submit any sensitive, secret or private recommendations before executing them. LLMs make mistakes.

First, we configure bia-bob:

```
# Load secret API key. You must unpack the contents of
# into the same folder before going ahead.
from dotenv import load_dotenv
load_dotenv()

import os
from bia_bob import bob
bob.initialize(endpoint='https://llm.scads.ai/v1', model='gpt-4o')
```

JupyterLab Interface:

- File Explorer: / docs / 02_bia_bob /
- Table of files:

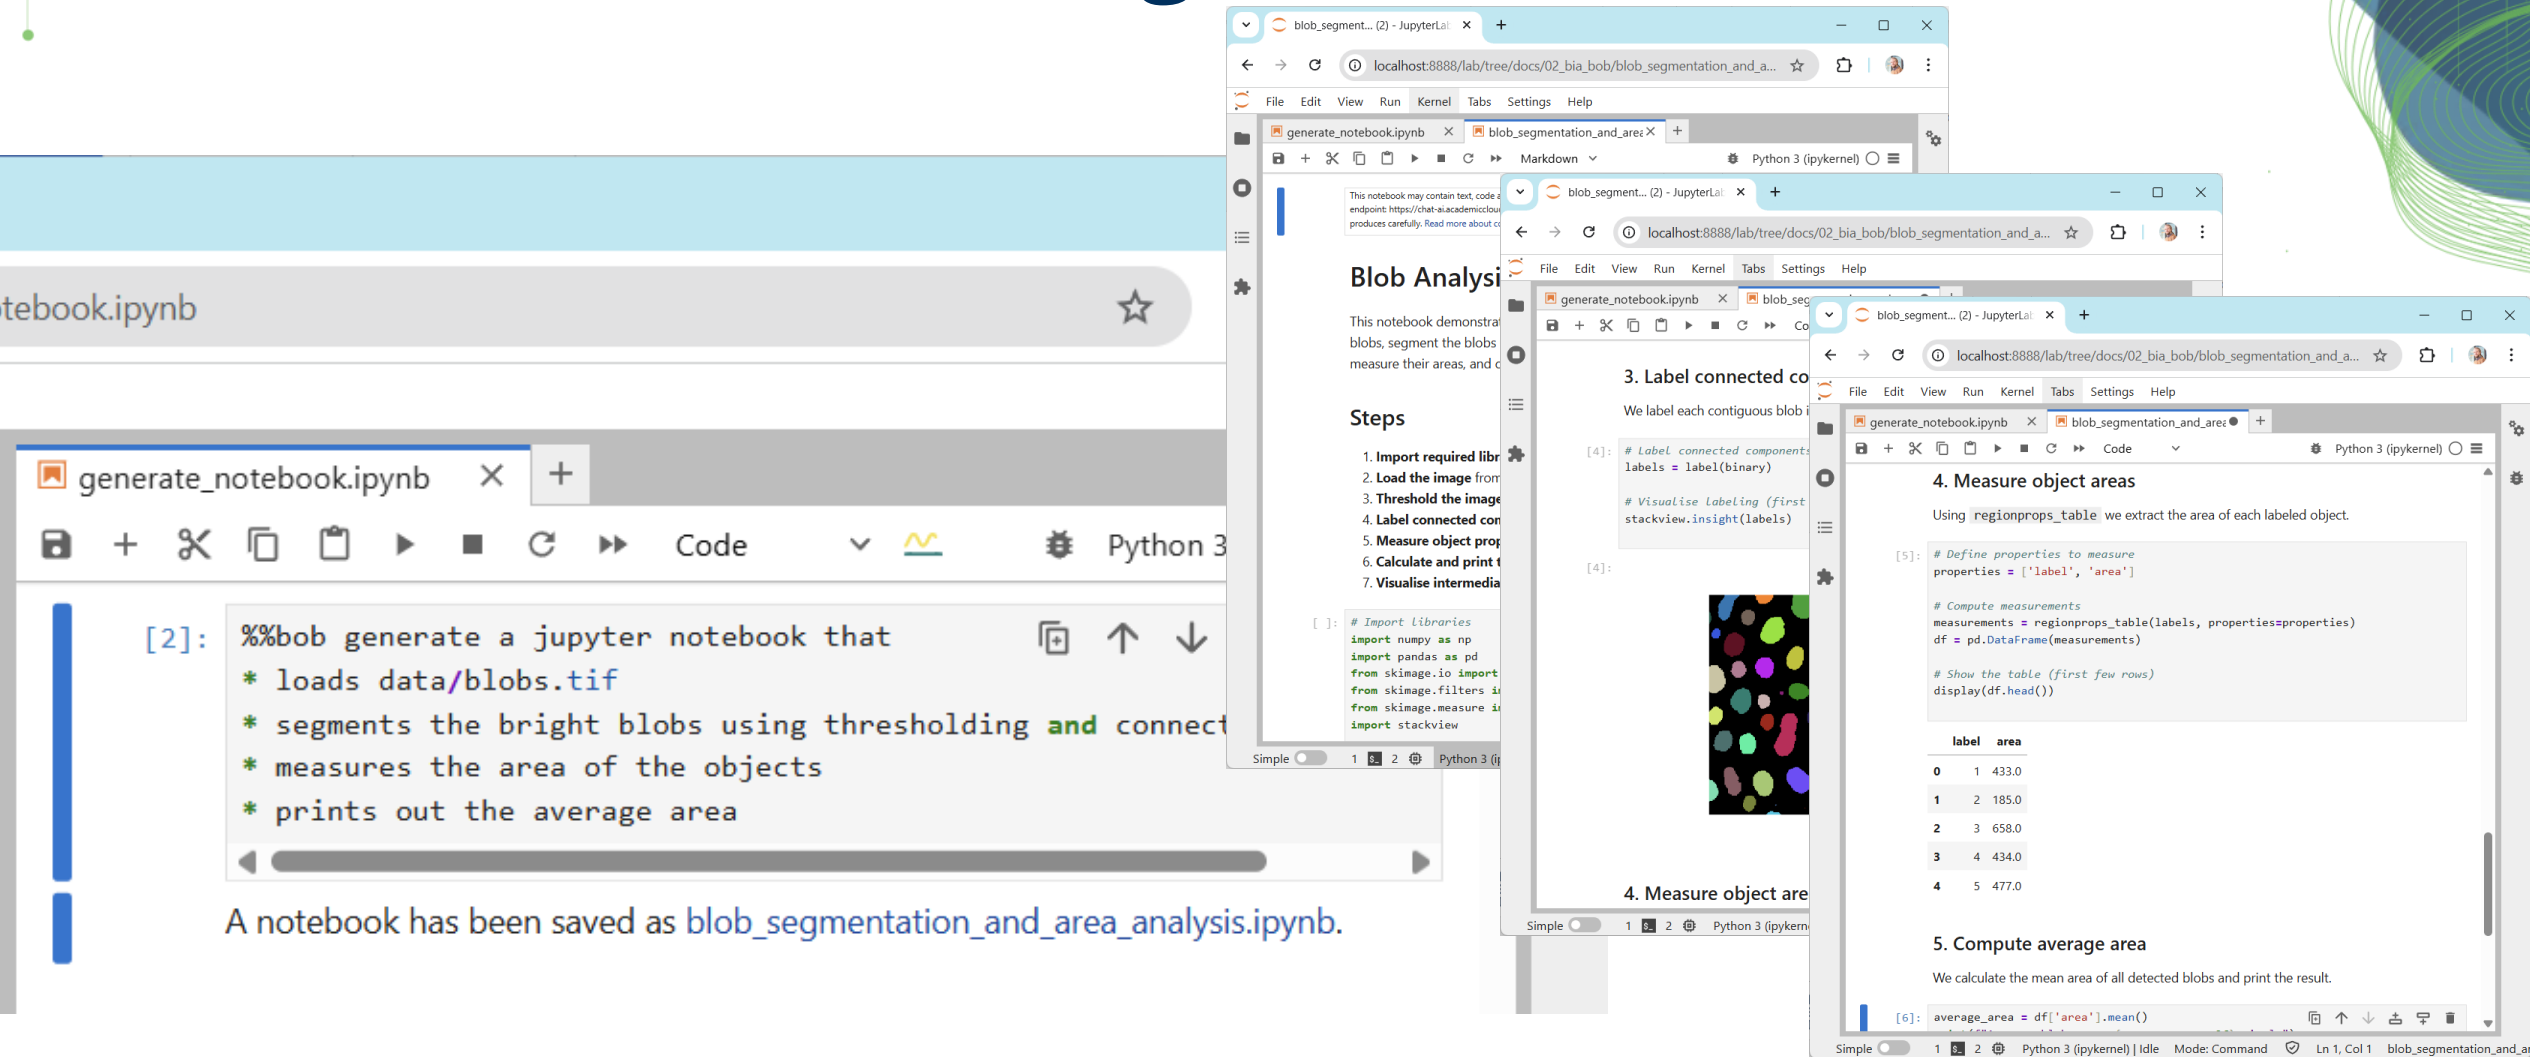
Name	Modified
data	1 hour ago
api_key.zip	21 days ago
bia_bob-statistical_testing.ipynb	24 days ago
bia-bob-demo.ipynb	21 days ago
generate_notebook.ipynb	24 days ago
translating_notebooks.ipynb	21 days ago
under_the_hood.ipynb	24 days ago
Untitled.ipynb	28 days ago

Code Editor:

```
[3]: %%bob
load data/blobs.tif and show it
```

Initialisation

Bia-Bob: notebook generation



The image shows a sequence of three JupyterLab notebook windows, illustrating the process of generating and running a notebook for blob segmentation and area analysis.

Window 1 (Left): Shows the initial notebook generation process. The code cell [2] contains the following instructions:

```
[2]: %%bob generate a jupyter notebook that
* loads data/blobs.tif
* segments the bright blobs using thresholding and connect
* measures the area of the objects
* prints out the average area
```

A message at the bottom states: "A notebook has been saved as blob_segmentation_and_area_analysis.ipynb."

Window 2 (Middle): Shows the generated notebook, "blob_segmentation_and_area_analysis.ipynb". It includes a table of contents and a list of steps:

1. Import required libraries
2. Load the image from data/blobs.tif
3. Threshold the image
4. Label connected components
5. Measure object properties
6. Calculate and print the average area
7. Visualise intermediate results

The code cell [4] shows the import of libraries:

```
[4]: # Import libraries
import numpy as np
import pandas as pd
from skimage.io import imread
from skimage.filters import threshold_otsu
from skimage.measure import label, regionprops
import stackview
```

A visualization of the segmented blobs is shown.

Window 3 (Right): Shows the execution of the notebook. The code cell [5] defines properties to measure and computes the area of each labeled object:

```
[5]: # Define properties to measure
properties = ['label', 'area']

# Compute measurements
measurements = regionprops_table(labels, properties=properties)
df = pd.DataFrame(measurements)

# Show the table (first few rows)
display(df.head())
```

The resulting table is displayed:

	label	area
0	1	433.0
1	2	185.0
2	3	658.0
3	4	434.0
4	5	477.0

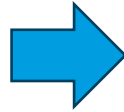
The code cell [6] calculates the average area:

```
[6]: average_area = df['area'].mean()
```

Code for statistical analysis

Be careful when prompting for code and **provide** as many **details** as you can!

%%bob load data/data.csv and determine if the columns A and B are significantly different.



```
import pandas as pd
import numpy as np
from scipy.stats import ttest_ind
import matplotlib.pyplot as plt

# 1. Load the CSV file
df = pd.read_csv('data/data.csv')

# 2. Keep only rows where both A and B are present
df = df[['A', 'B']].dropna()

# 3. Extract the two series
a = df['A'].values
b = df['B'].values

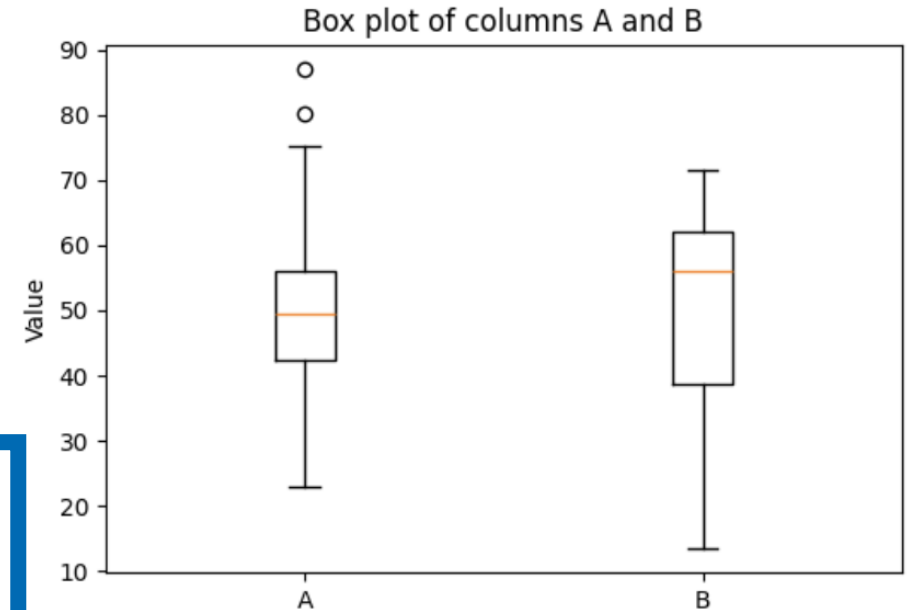
# 4. Perform independent two-sample t-test
t_stat, p_val = ttest_ind(a, b, equal_var=False) # Welch's t-test

# 5. Output the results
print(f"t-statistic: {t_stat:.4f}")
print(f"p-value: {p_val:.4e}")
if p_val < 0.05:
    print("=> The difference between columns A and B is statistically significant (α=0.05).")
else:
    print("=> No statistically significant difference detected (α=0.05).")

# 6. Visual comparison: box-plot
plt.figure(figsize=(6, 4))
plt.boxplot([a, b], labels=['A', 'B'])
plt.title('Box plot of columns A and B')
plt.ylabel('Value')
plt.show()
```



t-statistic: -1.3287
p-value: 1.8428e-01
=> No statistically significant difference detected ($\alpha=0.05$).



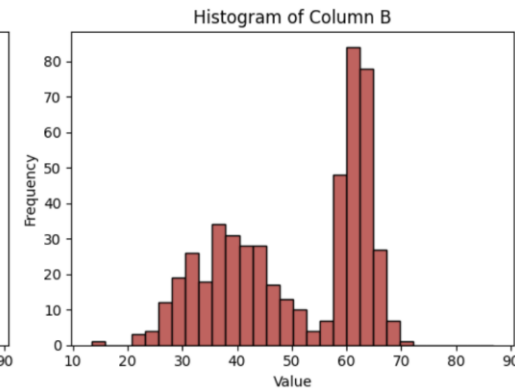
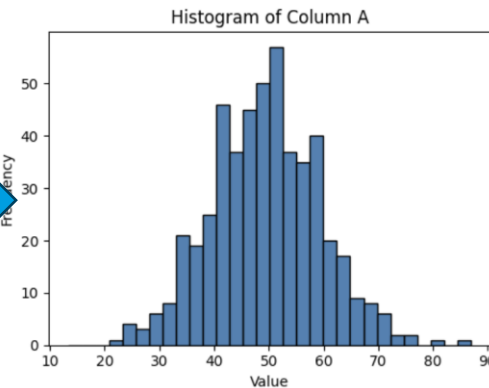
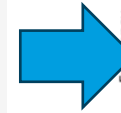
LLMs often provide “simple” solutions.

Code for statistical analysis

Be careful when prompting for code and **provide** as many **details** as you can!

%bob
Plot histograms **for A and B**

```
# Assuming df, a, b are already defined from the previous step.  
# If not, uncomment the following lines to load them again:  
# import pandas as pd  
# df = pd.read_csv('data/data.csv')[['A', 'B']].dropna()  
# a = df['A'].values  
# b = df['B'].values  
  
# Determine a common bin range for both histograms  
min_val = min(a.min(), b.min())  
max_val = max(a.max(), b.max())  
bins = 30 # you can adjust the number of bins  
  
plt.figure(figsize=(10, 4))  
  
# Histogram for column A  
plt.subplot(1, 2, 1)  
plt.hist(a, bins=bins, range=(min_val, max_val), color='steelblue', edgecolor='black')  
plt.title('Histogram of Column A')  
plt.xlabel('Value')  
plt.ylabel('Frequency')  
  
# Histogram for column B  
plt.subplot(1, 2, 2)  
plt.hist(b, bins=bins, range=(min_val, max_val), color='indianred', edgecolor='black')  
plt.title('Histogram of Column B')  
plt.xlabel('Value')  
plt.ylabel('Frequency')  
  
plt.tight_layout()  
plt.show()
```



Look at your
data!

Code for statistical analysis

Be careful when prompting for code and **provide** as many

`%%bob`

Please test again **if** columns A **and** B are significantly different.
Before applying a statistical test, ensure that the necessary pre-conditions are met.

```
from scipy.stats import shapiro, levene, mannwhitneyu
# Note: pandas (pd), numpy (np), and test_ind are already available.

# 1. Load data (keep original df/a/b untouched)
data_path = "data/data.csv"
df_raw = pd.read_csv(data_path)
# Keep only rows where both A and B are present
df_clean = df_raw[['A', 'B']].dropna()
a_vals = df_clean['A'].values
b_vals = df_clean['B'].values

# 2. Normality tests (Shapiro-Wilk)
shapiro_a_stat, shapiro_a_p = shapiro(a_vals)
shapiro_b_stat, shapiro_b_p = shapiro(b_vals)

# 3. Equality of variances (Levene)
lev_stat, lev_p = levene(a_vals, b_vals)

# 4. Choose appropriate hypothesis test
if shapiro_a_p > 0.05 and shapiro_b_p > 0.05:
    # both distributions look normal
    if lev_p > 0.05:
        test_name = "Student's t-test (equal variances)"
        t_stat, p_val = test_ind(a_vals, b_vals, equal_var=True)
    else:
        test_name = "Welch's t-test (unequal variances)"
        t_stat, p_val = test_ind(a_vals, b_vals, equal_var=False)
else:
    # at least one non-normal - use non-parametric test
    test_name = "Mann-Whitney U test (non-parametric)"
    # two-sided test
    t_stat, p_val = mannwhitneyu(a_vals, b_vals, alternatives='two-sided')

# 5. Print report
print("=== Assumption checks ===")
print(f"Shapiro-Wilk A: statistic={shapiro_a_stat:.4f}, p={shapiro_a_p:.4e}")
print(f"Shapiro-Wilk B: statistic={shapiro_b_stat:.4f}, p={shapiro_b_p:.4e}")
print(f"Levene variance test: statistic={lev_stat:.4f}, p={lev_p:.4e}")
print(f"Final hypothesis test used: {test_name}")
print(f"Statistic = {t_stat:.4f}")
print(f"P-value = {p_val:.4e}")
if p_val < 0.05:
    print("> Significant difference between columns A and B (alpha=0.05).")
else:
    print("> No significant difference detected between columns A and B (alpha=0.05).")
```

=== Assumption checks ===

Shapiro-Wilk A: statistic=0.9972, p=5.6236e-01 normal

Shapiro-Wilk B: statistic=0.8937, p=4.0635e-18 non-normal

Levene variance test: statistic=80.6718, p=1.3016e-18 unequal variances

=== Final hypothesis test ===

Test used: Mann-Whitney U test (non-parametric)

Statistic = 113657.0000

P-value = 1.3000e-02

=> Significant difference between columns A and B ($\alpha=0.05$).

LLMs cannot be *brave* scientists yet, but you can!

Acknowledgements

Collaborators & contributors

Christian Tischer (EMBL Heidelberg), Jean-Karim Hériché (EMBL Heidelberg), Nico Scherf (MPI CBS Leipzig), Loïc A. Royer (CZ Biohub), Elena Nicolay (UFZ Leipzig), Kevin Yamauchi (ETH Zurich / Basel), Seth Hinz, Teun A.P.M. Huijben, Jordão Bragantini (CZ Biohub San Francisco), Ilan Theodoro (CZ Biohub and the State University of Campinas, Brazil), Ian C. (University of British Columbia), Pradeep Rajasekhar (WEHI Australia)

BIDS Team @ ScaDS.AI / UL



Lea Kabjesz, Lea Gihlein, Mara Lampert

Communities & platforms



NFDI 4
BIOIMAGE



image.sc

Funding



Bundesministerium
für Forschung, Technologie
und Raumfahrt



SACHSEN Diese Maßnahme wird gefördert durch die Bundesregierung aufgrund eines Beschlusses des Deutschen Bundestages. Diese Maßnahme wird mitfinanziert durch Steuermittel auf der Grundlage des von den Abgeordneten des Sächsischen Landtags beschlossenen Haushaltes.



Deutsche
Forschungsgemeinschaft



Using LLMs
Bio Image Analysis
Robert Haase
@haesleinhuepf
September 2025

RH acknowledges the financial support by the Federal Ministry of Education and Research of Germany and by Sächsische Staatsministerium für Wissenschaft, Kultur und Tourismus in the programme Center of Excellence for AI-research „Center for Scalable Data Analytics and Artificial Intelligence Dresden/Leipzig“, project identification number: ScaDS.AI



TECHNISCHE
UNIVERSITÄT
DRESDEN



UNIVERSITÄT
LEIPZIG