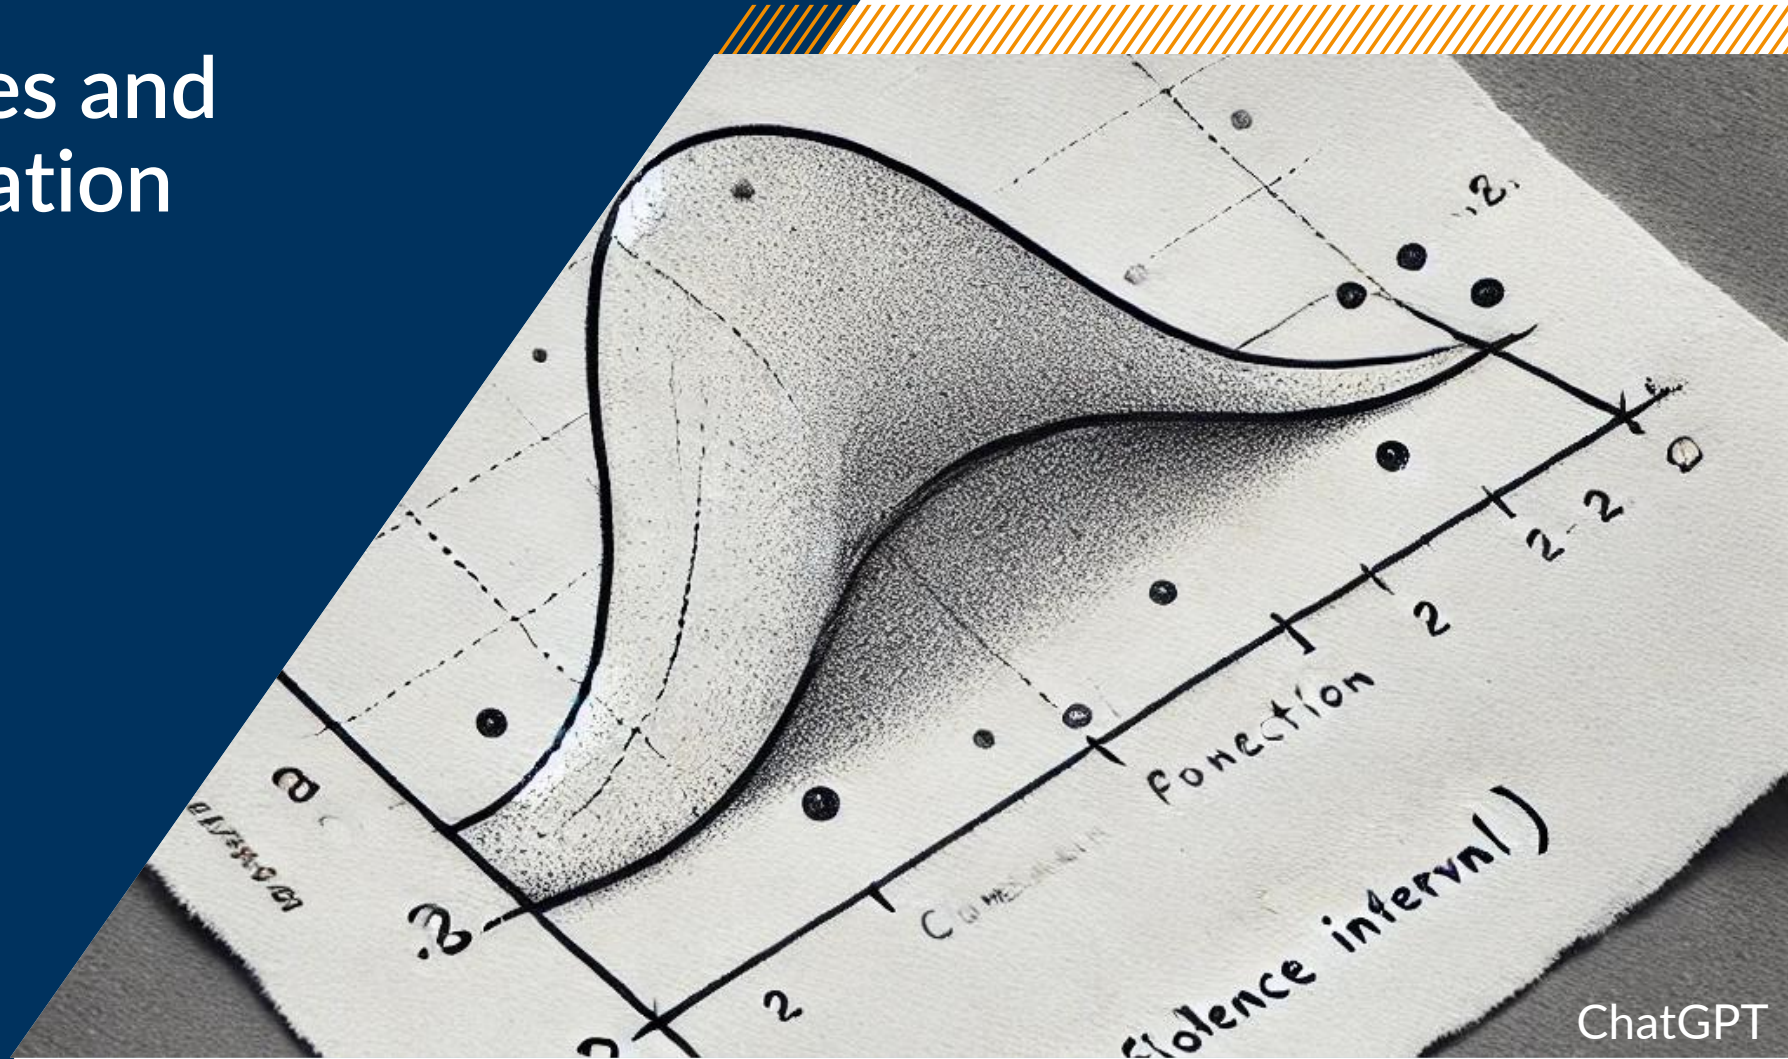


Gaussian Processes and Bayesian Optimization

- an introduction to core ideas
and practical implementation

Martin Rudolph

Leibniz Institute of Surface Engineering (IOM)
Modelling and Simulation



Agenda

- 0 Challenges in Process Optimization 11h START
- 1 Gaussian Process Intuition
- 2 GP Regression: from prior to posterior
 - Stochastic process
 - Bayesian conditioning / updating
 - 13h LUNCH
 - Log marginal likelihood
 - 16h COFFEE
- 3 Bayesian Optimization
 - 19h FESTIVAL OF LIGHTS

Opening comments

- Non-mathematical introduction to GP and BO
- Focus on using the GP/BO in scikit-learn for efficiently building response surfaces
 - Lab experiments
 - Computer simulations
 - Not included: hyperparameter optimization in ML
- Learn the concepts of GP and BO
 - Includes looking at equations, abbreviations are marked by
 - Point out the differences to other ML techniques
 - Practical sessions with (limited) code to illustrate the process in general
- Big group of students, one trainer
- Learning units: Lecture + Code examples + Practical session + Discussion





0 Challenges in process optimization



Challenges in process optimization



delonghi.de

Challenges in process optimization

- large number of possible process parameter combinations
 - grain size
 - water temperature
 - amount of powder
 - brew duration
 - weather (!!)
 - ...

5 parameter with 10 variations each:
 $10^5 = 100,000$ possible combinations
→ „combinatorial explosion“



delonghi.de

Challenges in process optimization

- large number of possible process parameter combinations
 - grain size
 - water temperature
 - amount of powder
 - brew duration
 - weather (!!)
 - ...


5 parameter with 10 variations each:
 $10^5 = 100,000$ possible combinations
→ „combinatorial explosion“



delonghi.de

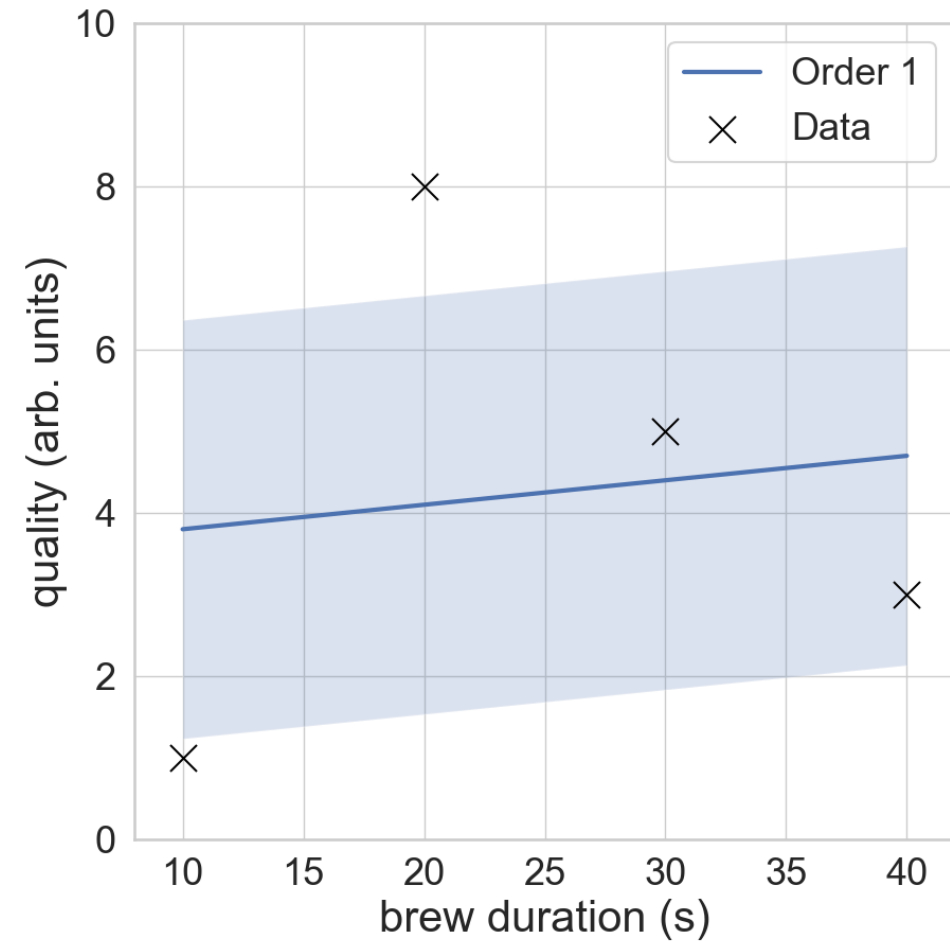
Response surfaces using polynomials

Single parameter: brew duration




Response surfaces using polynomials

Single parameter: brew duration




Response surfaces using polynomials

Single parameter: brew duration



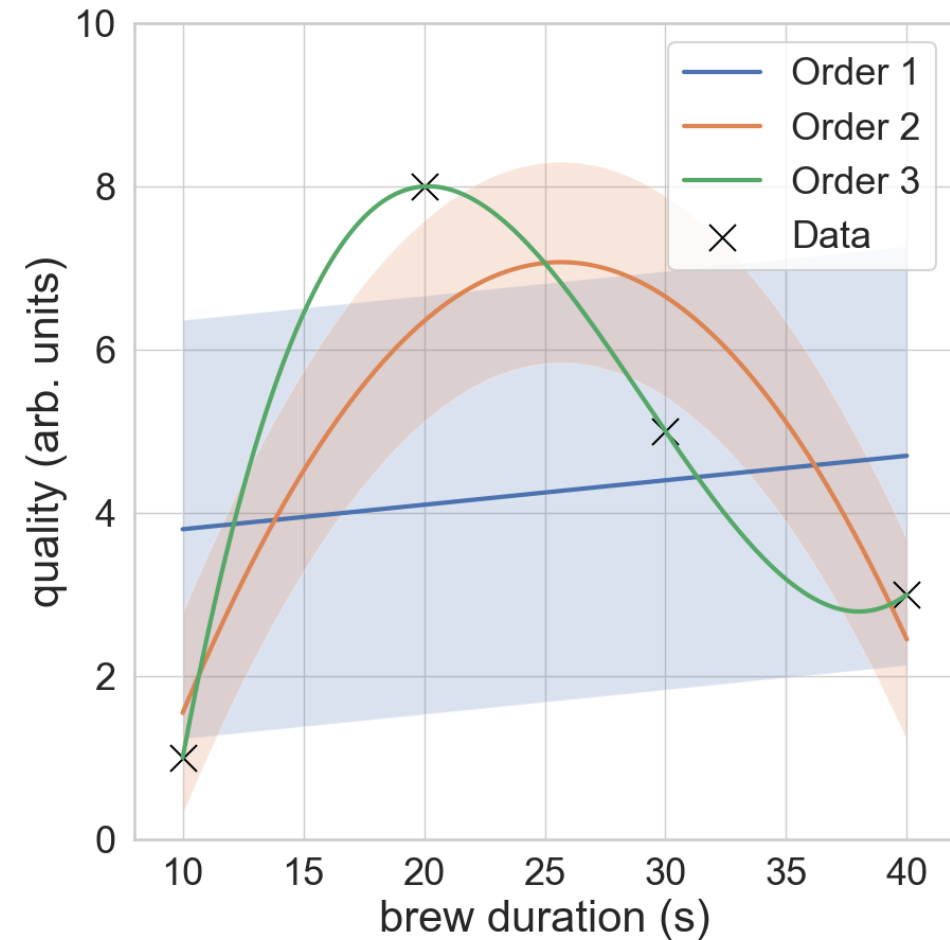
Response surfaces using polynomials

Single parameter: brew duration



Response surfaces using polynomials

- Single parameter: brew duration
- Limitations:
 - Need to choose a (rigid) model
 - Constant uncertainty (even in the vicinity of data points)
 - Perfect fit (no uncertainty) for 3rd order polynomial




Response surfaces using polynomials



- Single parameter: brew duration
- Limitations:
 - Need to choose a (rigid) model
 - Constant uncertainty (even in the vicinity of data points)
 - Perfect fit (no uncertainty) for 3rd order polynomial

Wishlist:

- No need to choose a rigid model
- Uncertainty measure that lives up to its name





1 Looking at a Gaussian Process



Danie Krige

Danie Krige
1919-2013



*Minitt et al. Trans. Roy. Soc. South Africa, 68, 2013.

1938 B.Sc. Mining Engineering

1938 Anglovaal: sampling and valuation of gold in the ground

*„decisions on new gold mines of critical importance to the State [...] were being taken on a limited number of drillholes, **without any scientific analysis** of the risks of failure“ **


1943 Government Mining Engineer

*Krige, Danie G. (1951). "A statistical approach to some basic mine valuation problems on the Witwatersrand". J. of the Chem., Metal. and Mining Soc. of South Africa. 52 (6): 119–139.
(today cited 4600 times)*



1951 Anglovaal: application of „kriging“ of ore reserves

1981 Witwatersrand University: Professor of Mineral Economics



Gaussian Process in action





Gaussian Process in action





Gaussian Process in action





Gaussian Process in action



Gaussian Process in action

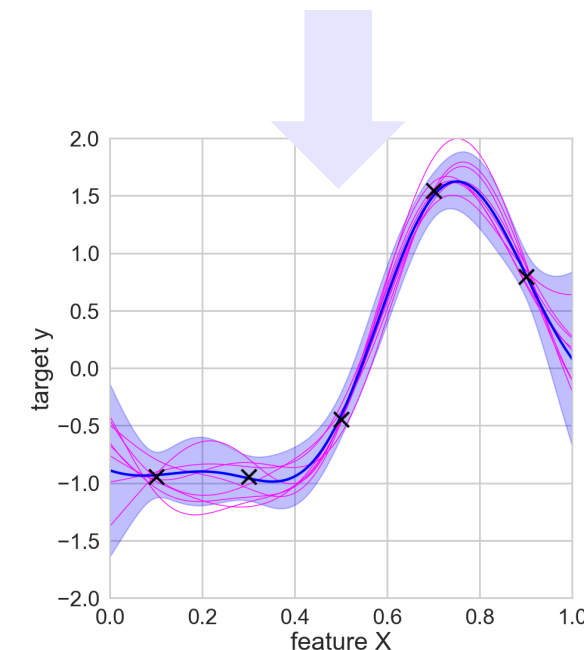



Gaussian Process in action



Advantages

- // Non-parametric model
 - // traditional parametric models: parameters define a specific functional shape
 - // e.g. linear fit $y = a_0 + a_1 x$
 - // (hyper-)parameters in GP define the *function space* (= infinite set of possible functions)
- // Adaptive
 - // add data points as experimental data comes in
 - // GP: model complexity can grow with more data points
- // Measure of uncertainty
 - // Collect data where uncertainty is high (exploration)
 - // Collect data where model prediction is best (exploitation)
 - // (→ Bayesian optimization, later today)



Questions?

Questions?

- How to generate the curves?
- Do these curves have a common property?



2 GP Regression

2.1 Stochastic Process


2.2 Bayesian Conditioning / Updating

2.3 Log marginal likelihood





Stochastic process

sampling from a (normal) distribution




mean μ
variance σ^2






Stochastic process

sampling from an **independent** multivariate normal distribution




...




mean vector μ
variance vector σ^2

Stochastic process


sampling from a **general** multivariate normal distribution



...




mean vector μ
covariance matrix σ^2



stochastic process based on
normal distributions: **Gaussian process**


Covariance matrix

- specifies covariance („correlation“) between each pair of random variables
 - $\text{Cov} = 1 \rightarrow$ strongest covariance
 - $\text{Cov} = 0 \rightarrow$ no covariance
- is a measure for distance between two points





Covariance matrix

- specifies covariance („correlation“) between each pair of random variables



Covariance matrix

- specifies covariance („correlation“) between each pair of random variables

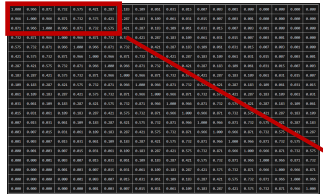


- Covariance function (or kernel function)
- Radial Basis Kernel


$$\text{cov}(f(x), f(x')) = k(x, x') = \exp\left(-\frac{(x-x')^2}{2l^2}\right)$$

Covariance matrix

- specifies covariance („correlation“) between each pair of random variables



	0	1	2	3	4	5	6	
0	1.000	0.966	0.871	0.732	0.575	0.421	0.287	
1	0.966	1.000	0.966	0.871	0.732	0.575	0.421	...
2	0.871	0.966	1.000	0.966	0.871	0.732	0.575	
				...				


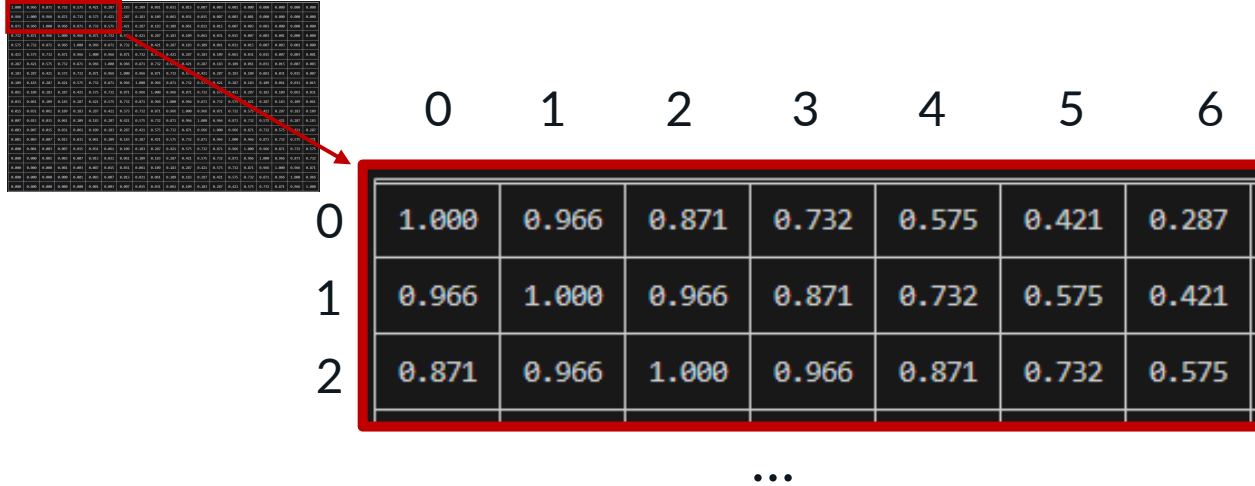


- Covariance function (or kernel function)
- Radial Basis Kernel (most common kernel function, but there are many more!)

$$\text{cov}(f(x), f(x')) = k(x, x') = \exp\left(-\frac{(x-x')^2}{2l^2}\right)$$

Covariance matrix

- specifies covariance („correlation“) between each pair of random variables



- Covariance function (or kernel function)
- Radial Basis Kernel (most common kernel function, but there are many more!)


$$\text{cov}(f(x), f(x')) = k(x, x') = \exp\left(-\frac{(x-x')^2}{2l^2}\right)$$

length scale parameter l

Length scale parameter

$l = 4$



	0	1	2	3	4	5	6	
0	1.000	0.966	0.871	0.732	0.575	0.421	0.287	
1	0.966	1.000	0.966	0.871	0.732	0.575	0.421	...
2	0.871	0.966	1.000	0.966	0.871	0.732	0.575	
	...							



Length scale parameter

$l = 4$

	0	1	2	3	4	5	6	
0	1.000	0.966	0.871	0.732	0.575	0.421	0.287	
1	0.966	1.000	0.966	0.871	0.732	0.575	0.421	...
2	0.871	0.966	1.000	0.966	0.871	0.732	0.575	
			...					







Length scale parameter

$l = 2$

	0	1	2	3	4	5	6	
0	1.000	0.871	0.575	0.287	0.109	0.031	0.007	
1	0.871	1.000	0.871	0.575	0.287	0.109	0.031	...
2	0.575	0.871	1.000	0.871	0.575	0.287	0.109	
	...							

$l = 4$

	0	1	2	3	4	5	6	
0	1.000	0.966	0.871	0.732	0.575	0.421	0.287	
1	0.966	1.000	0.966	0.871	0.732	0.575	0.421	...
2	0.871	0.966	1.000	0.966	0.871	0.732	0.575	
	...							



Length scale parameter

$l = 2$







	0	1	2	3	4	5	6	
0	1.000	0.871	0.575	0.287	0.109	0.031	0.007	
1	0.871	1.000	0.871	0.575	0.287	0.109	0.031	...
2	0.575	0.871	1.000	0.871	0.575	0.287	0.109	
	...							

$l = 4$

	0	1	2	3	4	5	6	
0	1.000	0.966	0.871	0.732	0.575	0.421	0.287	
1	0.966	1.000	0.966	0.871	0.732	0.575	0.421	...
2	0.871	0.966	1.000	0.966	0.871	0.732	0.575	
	...							

$l = 8$

	0	1	2	3	4	5	6	
0	1.000	0.991	0.966	0.925	0.871	0.805	0.732	
1	0.991	1.000	0.991	0.966	0.925	0.871	0.805	...
2	0.966	0.991	1.000	0.991	0.966	0.925	0.871	
	...							



Tasks

- Run 01_gaussian_process.py
- What does the plot show?
- Vary the length scale parameter. What do you observe?
- What do you observe when changing the RBF kernel to
 - Matern kernel
 - Periodic kernel (exp-sine-squared)
 - Linear kernel

Tasks

Run 01_gaussian_process.py

What does the plot show?

Vary the length scale parameter. What do you observe?

What do you observe when changing the RBF kernel to

Matern kernel

(becomes RBF kernel for $\nu = \infty$)

$$k(x_i, x_j) = \frac{1}{\Gamma(\nu)2^{\nu-1}} \left(\frac{\sqrt{2\nu}}{l} d(x_i, x_j) \right)^\nu K_\nu \left(\frac{\sqrt{2\nu}}{l} d(x_i, x_j) \right)$$

Periodic kernel (ExpSineSquared)

$$k(x_i, x_j) = \exp \left(-\frac{2 \sin^2(\pi d(x_i, x_j)/p)}{l^2} \right)$$

Linear kernel (DotProduct)

$$k(x_i, x_j) = \sigma_0^2 + x_i \cdot x_j$$





2 GP Regression


2.1 Stochastic Process

2.2 **Conditioning / Updating the GP**


2.3 Model Assessment



Conditioning / Updating the GP




Conditioning / Updating the GP



Very intuitive picture, but in reality one works with the underlying probability distributions.

GP Prior (before any data)


- Prior: probability distribution over functions before adding observations



infinite number of normal distributions with
a **mean vector** and a **covariance matrix**

GP Prior (before any data)


- Prior: probability distribution over functions before adding observations



infinite number of normal distributions with
a **mean vector** and a **covariance matrix**

GP Prior (before any data)


- Prior: probability distribution over functions before adding observations



infinite number of normal distributions with
a **mean vector** and a **covariance matrix**

GP Prior (before any data)


- Prior: probability distribution over functions before adding observations



infinite number of normal distributions with
a **mean vector** and a **covariance matrix**

GP Prior (before any data)

- ▀ Prior: probability distribution over functions before adding observations

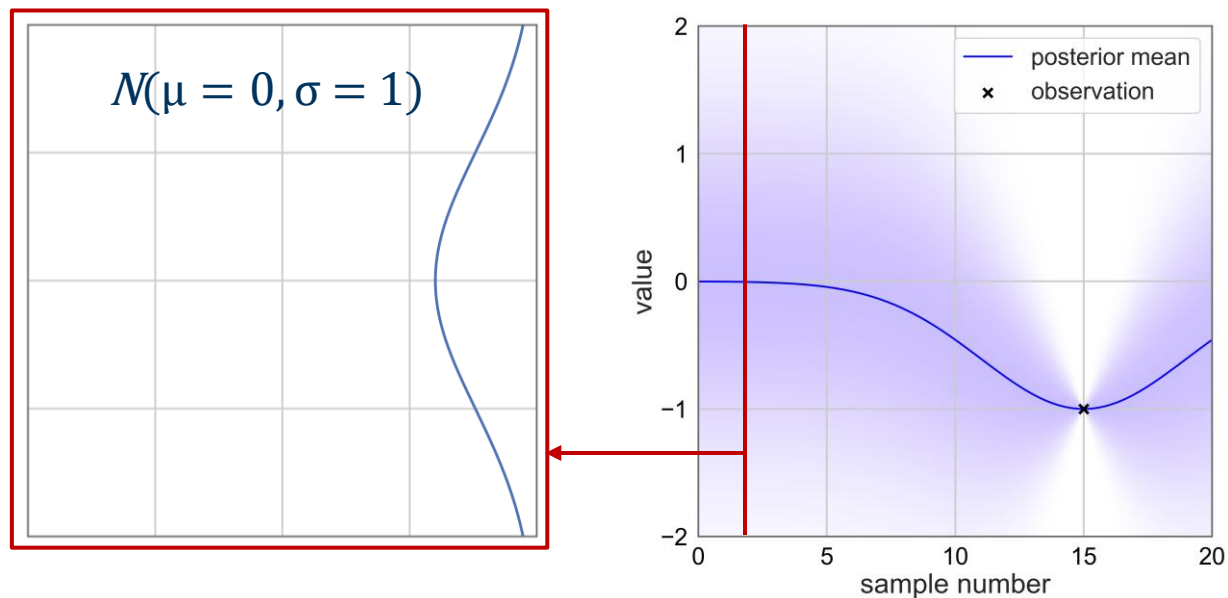


Conditioning: $\text{Posterior} \propto \text{Prior} \times \text{Likelihood}$



GP Posterior (after adding data)


- ▀ Prior: probability distribution over functions before adding observations



Conditioning: Posterior \propto Prior \times Likelihood

GP Posterior (after adding data)

- ▀ Prior: probability distribution over functions before adding observations




Conditioning: $\text{Posterior} \propto \text{Prior} \times \text{Likelihood}$

Likelihood



Conditioning: Posterior \propto Prior \times Likelihood

Likelihood: normal distribution with μ and sigma

Prior
 $N(\mu = 0, \sigma = 1)$




\times Likelihood \propto
 $N(\mu = -1, \sigma = 0.01)$



„slice through posterior“

$N(\mu = -0.99, \sigma = 0.01)$




Likelihood


Conditioning: Posterior \propto Prior \times Likelihood

Likelihood: normal distribution with μ and sigma


Prior
 $N(\mu = 0, \sigma = 1)$



Likelihood \propto
 $N(\mu = -1, \sigma = 0.01)$




Posterior






„slice through posterior“


$N(\mu = -0.99, \sigma = 0.01)$



$N(\mu = -1, \sigma = 0.3)$



$N(\mu = -0.91, \sigma = 0.29)$



Tasks

- Open 02_bayesian_condition.py
- What does the plot show?
- Reduce the length scale. What do you observe?
- Increase the length scale. What do you observe?
- Add additional data points. What do you observe?

Tasks

- Open 02_bayesian_condition.py
- What does the plot show?
- Reduce the length scale. What do you observe?
- Increase the length scale. What do you observe?
 - numerical instabilities due to ill-conditioned matrix
- Add additional data points. What do you observe?



2 GP Regression

2.1 Stochastic Process

2.2 Conditioning / Updating the GP

2.3 Model Assessment




Model assessment

objective: choose hyperparameters for


- // goodness of fit
 - // generalization
 - // in classical (deterministic) ML methods (support vector machines, random forrests,...):
 - // hyperparameters are properties of the method
 - // need to be adjusted to
 - // predict test data (**generalization**)
 - // with low error (**goodness of fit**)
- } Cross-validation to balance generalization and goodness of fit
- // in **probabilistic ML methods (GP)**
 - // hyperparameter (length scale in a GP) defines the probability distribution, that can be regarded as generating the data.
 - // (in practice: One can still do cross-validation to check the generalization)

Relation between lengthscale of prior and data


$l = 0.03$



$l = 0.17$




$l = 0.84$




Relation between lengthscale of prior and data


$l = 0.03$



$l = 0.17$




$l = 0.84$




Relation between lengthscale of prior and data


$l = 0.03$



$l = 0.17$




$l = 0.84$




Relation between lengthscale of prior and data


$l = 0.03$



$l = 0.17$




$l = 0.84$




Relation between lengthscale of prior and data


$l = 0.03$
„wiggly“ samples



$l = 0.17$
optimum length scale




$l = 0.84$
„flat“ samples




Relation between lengthscale of prior and data


$l = 0.03$
„wiggly“ samples



$l = 0.17$
optimum length scale



$l = 0.84$
„flat“ samples



Increased goodness of fit

Reduced model complexity (Volume of function space \rightarrow determinant of covariance matrix $|K|$)

\rightarrow Possibility to select an optimum lengthscale from the knowledge of the data!

Log likelihood

Likelihood function measures how probable it is to observe data \mathbf{y} at input \mathbf{X} given the hyperparameters θ .



to be maximized

$$\log p(\mathbf{y}|\mathbf{X}, \theta) = \underbrace{-\frac{1}{2} \mathbf{y}^T K_y^{-1} \mathbf{y}}_{\text{data fit term}} \quad \underbrace{-\frac{1}{2} \log |K_y|}_{\text{complexity penalty}} \quad \underbrace{-\frac{n}{2} \log 2\pi}_{\text{normalization (= constant)}}$$

Log likelihood

Likelihood function measures how probable it is to observe data \mathbf{y} at input \mathbf{X} given the hyperparameters θ .




to be maximized

$$\log p(\mathbf{y}|\mathbf{X}, \theta) = \underbrace{-\frac{1}{2} \mathbf{y}^T K_y^{-1} \mathbf{y}}_{\text{data fit term}} \quad \underbrace{-\frac{1}{2} \log |K_y|}_{\text{complexity penalty}} \quad \underbrace{-\frac{n}{2} \log 2\pi}_{\text{normalization (= constant)}}$$

Mahalanobis distance² =
generalized squared error
(more negative for worse fits)

determinant $|K_y|$
becomes larger for
small length scales (vectors
become linearly independent)

Log likelihood



Likelihood function measures how probable it is to observe data \mathbf{y} at input \mathbf{X} given the hyperparameters θ .

$$\log p(\mathbf{y}|\mathbf{X}, \theta) = \underbrace{-\frac{1}{2} \mathbf{y}^T \mathbf{K}_y^{-1} \mathbf{y}}_{\text{data fit term}} \quad \underbrace{-\frac{1}{2} \log |\mathbf{K}_y|}_{\text{complexity penalty}} \quad \underbrace{-\frac{n}{2} \log 2\pi}_{\text{normalization (= constant)}}$$

Tasks

- // Open 03_lengthscales_optimization.py
- // Run the script and optimize the lengthscales manually.
- // Have the lengthscales optimized automatically
 - // Set attribute `optimizer = „fmin_l_bfgs_b”`
- // Print the optimum length scale to the screen and compare to your lengthscales.
 - // Kernel can be accessed through `gp.kernel_`
 - // Lengthscale can be accessed through `gp.kernel_.length_scale`
- // Optimize the noise as well
 - // Use a combination of a *WhiteKernel* and an *RBF* kernel (composite kernel)
 - // Noise can be accessed through `gp.kernel_.k1.length_scale`
`gp.kernel_.k2.noise_level`

Tasks

- // Open 03_lengthscales_optimization.py

- // Run the script and optimize the lengthscale manually.

- // Have the lengthscale optimized automatically
 - // Set attribute `optimizer = „fmin_l_bfgs_b”`
 - // Possibly set the bounds `kernel = RBF(length_scale=1.0, length_scale_bounds=(1e-2, 1))`
 - // Possibly set restarts `gp = GaussianProcessRegressor(... , n_restarts_optimizer=10)`

- // Print the optimum length scale to the screen and compare to your lengthscale.
 - // Kernel can be accessed through `gp.kernel_`
 - // Lengthscale can be accessed through `gp.kernel_.length_scale`

- // Optimize the noise as well
 - // Use a combination of a *WhiteKernel* and an *RBF* kernel (composite kernel)
 - // Noise can be accessed through `gp.kernel_.k1.length_scale`
`gp.kernel_.k2.noise_level`
 - // Possibly add data points ;)



3 Bayesian Optimization




Where to sample next?

Probabilistic models offer

- /// Predictive mean
- /// Uncertainty estimate

Two strategies

- /// **Exploitation**
 - /// sample at the models' **best prediction** (max/min. of the predictive mean)
- /// **Exploration**
 - /// Sample at model's highest uncertainty for **maximum model improvement**



Acquisition function

- acquisition function:
 - calculates a score for sampling at a location given the current state of the model
- Simple example: upper confidence bound
 - „pick the point with the largest optimistic estimate“

$$\alpha_{UCB}(x) = \mu(x) + \beta\sigma(x) \quad (\text{for maximization problems})$$


$$\alpha_{UCB}(x) = \mu(x) - \beta\sigma(x) \quad (\text{for minimization problems})$$



exploration parameter β :

$\beta > 1$: promoting exploration
 $\beta < 1$: promoting exploitation

} possibility to
adjust with each
data point



Acquistion function

- acquisition function:
 - calculates a score for sampling at a location given the current state of the model
- Simple example: upper confidence bound
 - „pick the point with the largest optimistic estimate“

$$\alpha_{UCB}(x) = \mu(x) + \beta\sigma(x) \quad (\text{for maximization problems})$$


$$\alpha_{UCB}(x) = \mu(x) - \beta\sigma(x) \quad (\text{for minization problems})$$



exploration parameter β :

$\beta > 1$: promoting exploration
 $\beta < 1$: promoting exploitation


} possibility to
adjust with each
data point




How does the UBC
acquisition function
look like for this plot?

Upper confidence bound ($\beta = 1$)


1



2








3

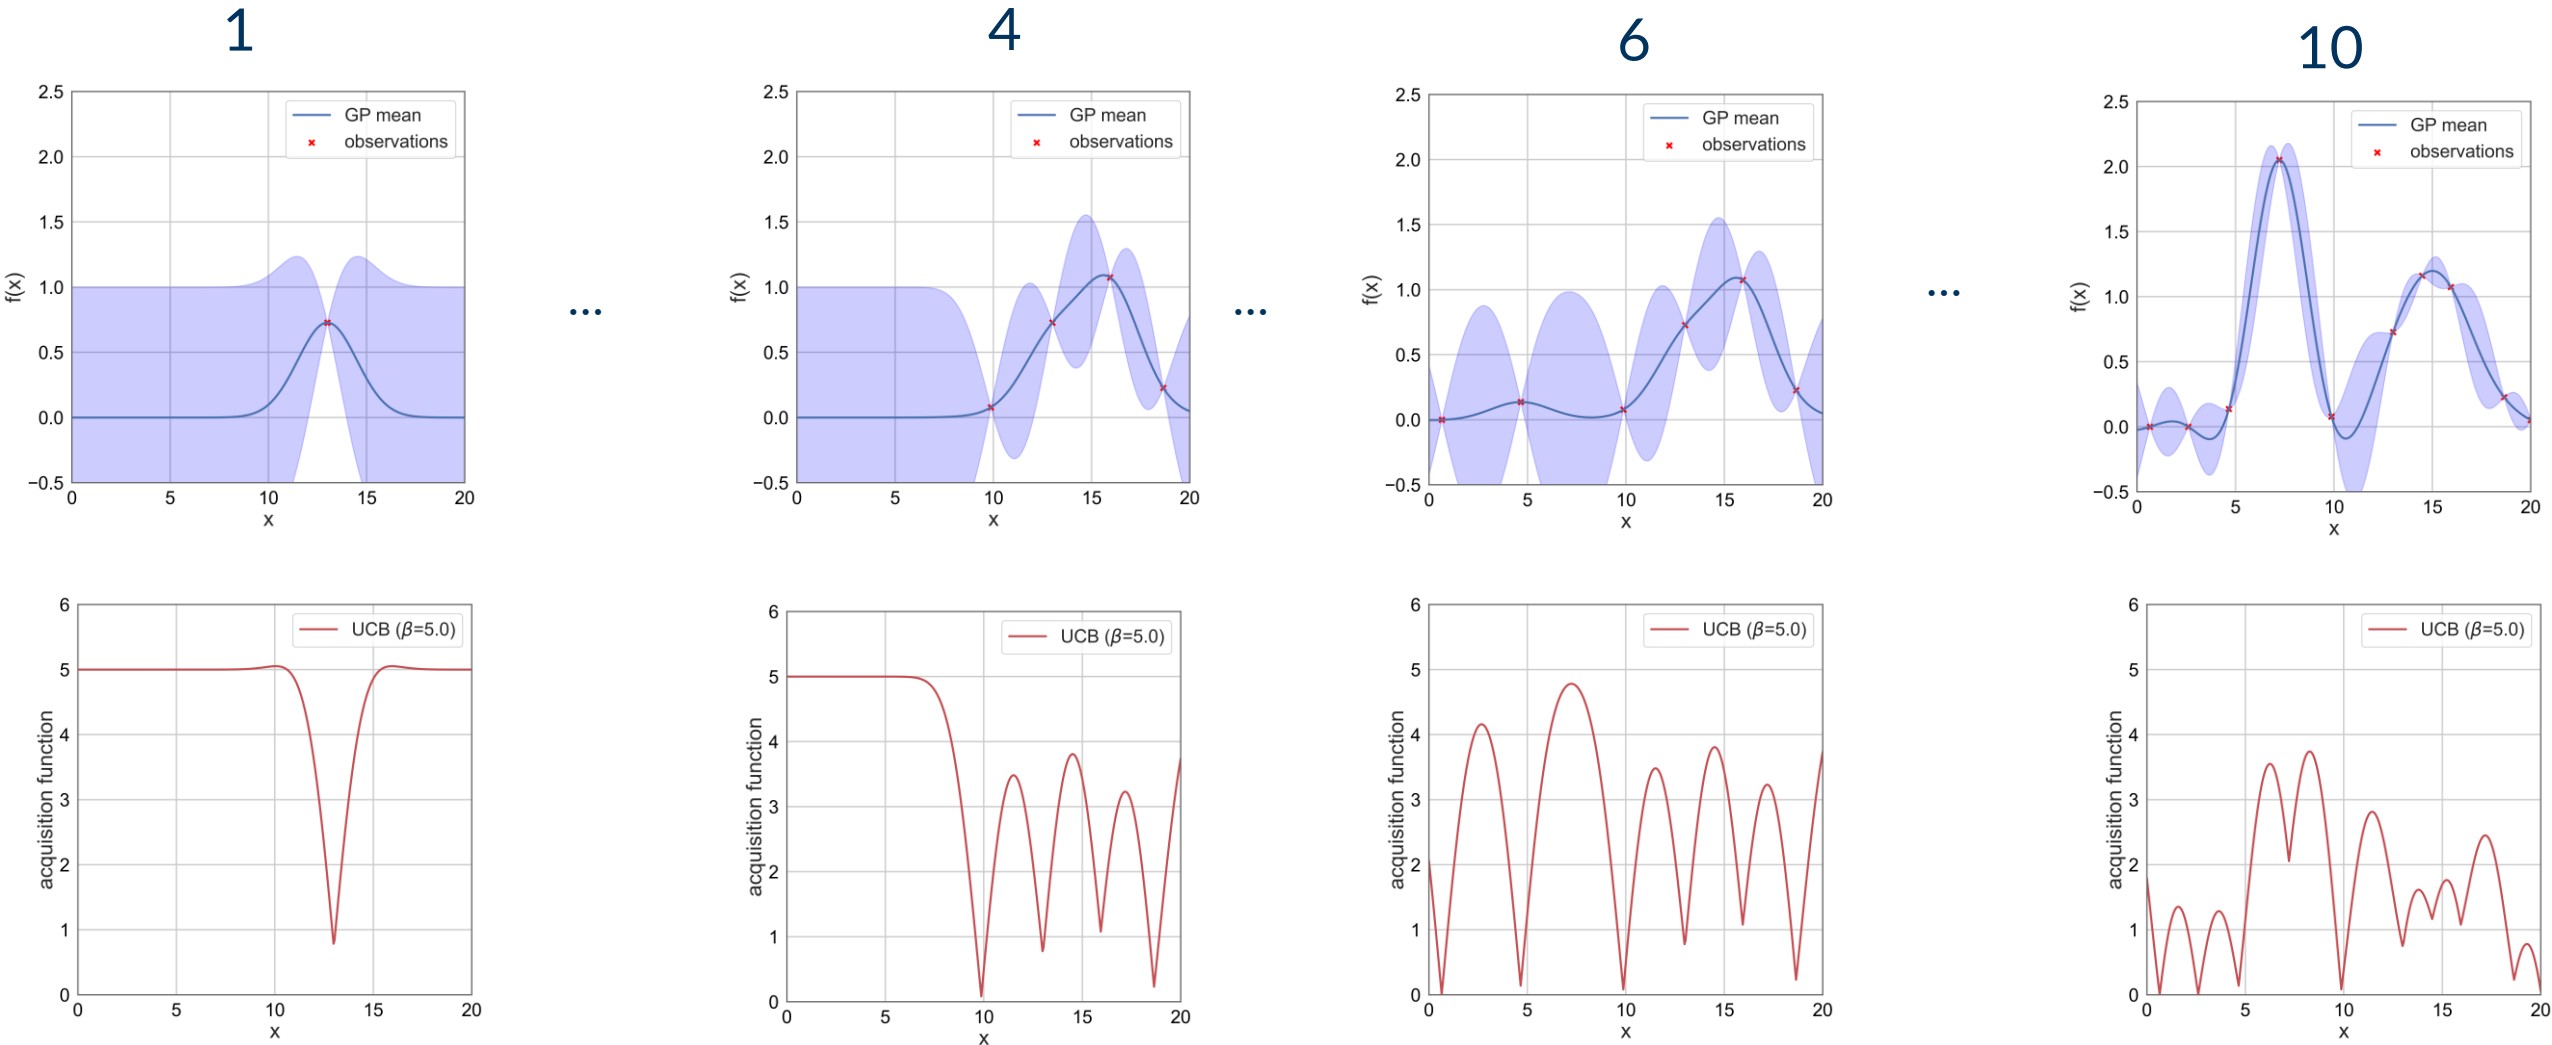


...


10



Upper confidence bound ($\beta = 5 \rightarrow$ promoting exploration)



Acquisition function II: Expected improvement



Expected improvement

$$\alpha_{EI}(x) = (\mu(x) - f(x^*)) \phi(z) + \sigma(x) \varphi(z)$$

$f(x^*)$ current best value

$$z(x) = \frac{\mu(x) - f(x^*)}{\sigma(x)}$$

$\phi(z)$ cumulative distribution function

$\varphi(z)$ probability distribution function

Acquisition function II: Expected improvement

Expected improvement

$$\alpha_{EI}(x) = \underbrace{(\mu(x) - f(x^*))}_{\text{expected gain of improvement}} \underbrace{\phi(z)}_{\text{chance of achieving it}} + \underbrace{\sigma(x)}_{\text{uncertainty}} \underbrace{\varphi(z)}_{\text{maximum, when } \mu(x) = f(x^*), \text{ and symmetric decrease in both directions}}$$

„exploitation term“

„exploration term“

$f(x^*)$ current best data point

$$z(x) = \frac{\mu(x) - f(x^*)}{\sigma(x)}$$

$\phi(z)$ cumulative distribution function

$\varphi(z)$ probability distribution function

Acquisition function II: Expected improvement

Expected improvement in practice

$$\alpha_{EI}(x) = \overbrace{(\mu(x) - f(x^*) - \xi) \phi(z)}^{\text{„exploitation term“}} + \overbrace{\sigma(x) \varphi(z)}^{\text{„exploration term“}}$$

$$z(x) = \frac{\mu(x) - f(x^*) - \xi}{\sigma(x)}$$


$\phi(z)$ cumulative distribution function

$\varphi(z)$ probability distribution function

Choose $\xi > 0$ for promoting exploration (e.g. 0.05 to 0.5)


Expected improvement

1




...

4








...

6



...

10




Tasks

- Open 04_bayesian_optimization.ipynb
- Approximate the black-box function
 - Do Gaussian Process Regression
 - Calculate the acquisition function (upper confidence bound)
 - Evaluate the black-box function and repeat
- Now use the expected improvement acquisition function

```
def expected_improvement(X, gp, y_best, xi=0.01):  
  
    mu, sigma = gp.predict(X, return_std=True)  
    sigma = sigma.reshape(-1, 1)  
    mu = mu.reshape(-1, 1)  
  
    # avoid division by zero  
    sigma = np.maximum(sigma, 1e-9)  
  
    imp = mu - y_best - xi  
    Z = imp / sigma  
  
    ei = imp * norm.cdf(Z) + sigma * norm.pdf(Z)  
    return ei.ravel()
```

Tasks

- Open 04_bayesian_optimization.ipynb
 - Approximate the black-box function
 - Do Gaussian Process Regression
 - Calculate the acquisition function (upper confidence bound)
 - Evaluate the black-box function and repeatif the optimization gets stuck, adjust β
 - Now use the expected improvement acquisition function
- if the optimization gets stuck, adjust
- ξ



```
def expected_improvement(X, gp, y_best, xi=0.01):  
  
    mu, sigma = gp.predict(X, return_std=True)  
    sigma = sigma.reshape(-1, 1)  
    mu = mu.reshape(-1, 1)  
  
    # avoid division by zero  
    sigma = np.maximum(sigma, 1e-9)  
  
    imp = mu - y_best - xi  
    Z = imp / sigma  
  
    ei = imp * norm.cdf(Z) + sigma * norm.pdf(Z)  
    return ei.ravel()
```



4 Final remarks



Summary

- // „Coffee machine“ challenge
 - // Numerical explosion
 - // Drawbacks with polynomial regression
- // Gaussian Process Regression
 - // Provides mean and uncertainty
 - // Non-parametric function
- // Co-variance matrix / kernel function
 - // Radial basis function kernel (RBF)
 - // Periodic kernel
 - // Linear kernel
- // Bayesian optimization / acquisition function
 - // Upper/lower confidence bound
 - // Expected improvement

Literature: Online docs for practioners

The screenshot shows the scikit-learn User Guide for Gaussian Processes. The left sidebar lists the navigation structure, with '1.7. Gaussian Processes' selected. The main content area is titled '1.7. Gaussian Processes' and describes them as a nonparametric supervised learning method. It lists advantages (interpolation, probabilistic prediction, versatility) and disadvantages (non-sparse, inefficient in high dimensions). A subsection '1.7.1. Gaussian Process Regression (GPR)' is also visible.

https://scikit-learn.org/stable/modules/gaussian_process.html

The screenshot shows the scikit-optimize documentation for Bayesian optimization with skopt. The left sidebar lists navigation options, with 'Bayesian optimization with skopt' selected. The main content area is titled 'Bayesian optimization with skopt' and includes a 'Problem statement' section. It describes the goal of finding the minimum of a function $f(x)$ under constraints. A 'Disclaimer' section notes that this is a better algorithm than Bayesian optimization if constraints are present. A 'Bayesian optimization loop' section describes the iterative process for $t = 1 : T$:


- Given observations $(x_i, y_i = f(x_i))$ for $i = 1 : t$, build a probabilistic model for the objective f . Integrate out all possible true functions, using Gaussian process regression.
- optimize a cheap acquisition/utility function u based on the posterior distribution for sampling the next point.
 $x_{t+1} = \arg \min_x u(x)$ Exploit uncertainty to balance exploration against exploitation.
- Sample the next observation y_{t+1} at x_{t+1} .

The bottom of the page shows the 'Expected improvement (default)' formula: $-EI(x) = -\mathbb{R}[f(x) - f(x^*)]$.


https://scikit-optimize.github.io/stable/auto_examples/bayesian-optimization.html

Books for further studies

Rasmussen and Williams: Gaussian Processes for Machine Learning, MIT Press, 2005.



Garnett: Bayesian Optimization, Cambridge University Press, 2023.



Questions